

---

# **Learning with Small Samples**

## **Including zero-shot learning**

Nour Karessli  
DSR 2019

---



# Nour Karessli

Data Scientist, Zalando SE  
Computer vision engineer, EyeEm GmbH

Master in computer science, Saarland University - Max Planck Institute  
Bachelor in software engineering, Damascus University

[nour.karessli@gmail.com](mailto:nour.karessli@gmail.com)  
[LinkedIn](#)

---

# About this tutorial

- A basic understanding of zero-shot and low-shot learning
- Get to know state of the art approaches
- Hands-on experience with zero-shot image classification
- Hands-on experience with training image classifier with small set

---

# Prerequisites

- Basic math e.g. matrix operations, derivatives,.., etc
- Basic understanding of ML concepts e.g. classifier, loss function,.., etc.
- Basic DL concepts e.g. MLP, CNN, LSTM, ..,etc
- Python, Keras

---

# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises

---

# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises

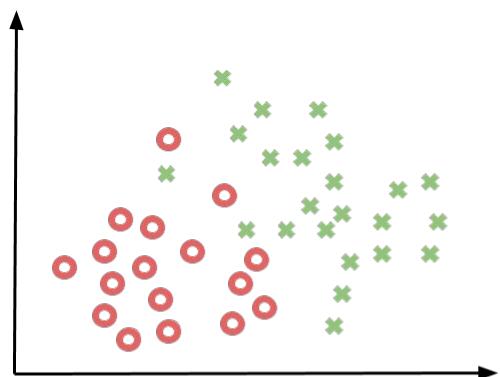


# Introduction & Motivation



# Learning

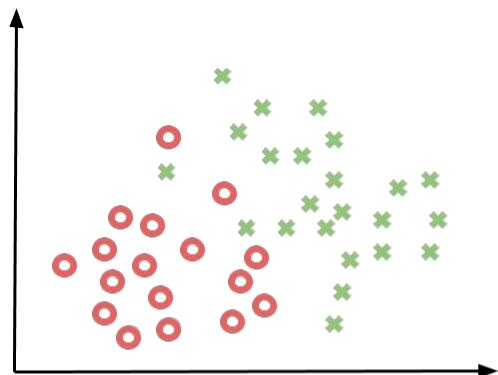
Supervised



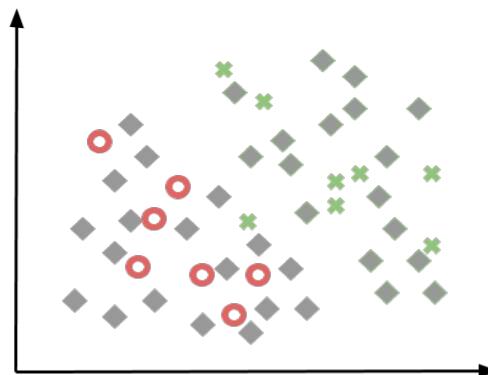
---

# Learning

Supervised



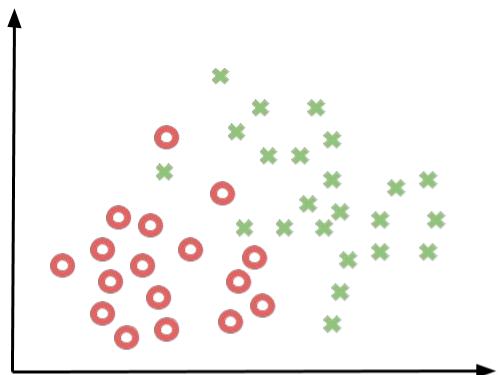
Semi-supervised



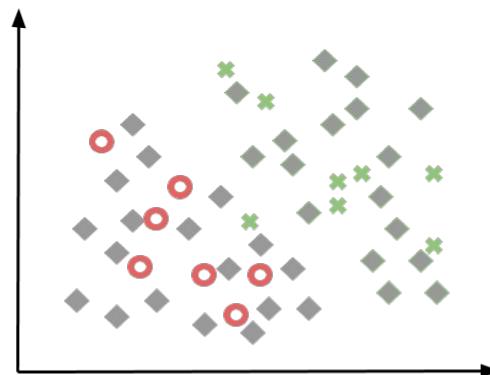
---

# Learning

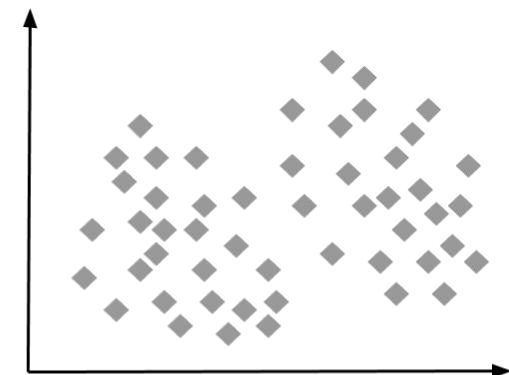
Supervised



Semi-supervised



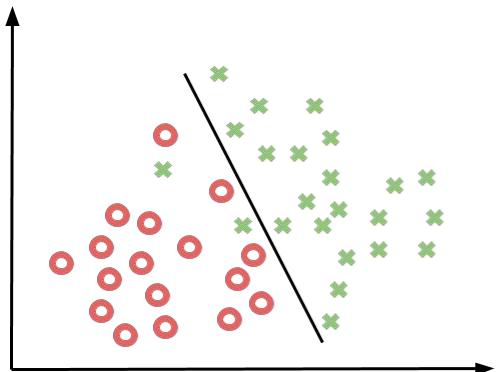
Unsupervised



---

# Learning

Supervised



Semi-supervised



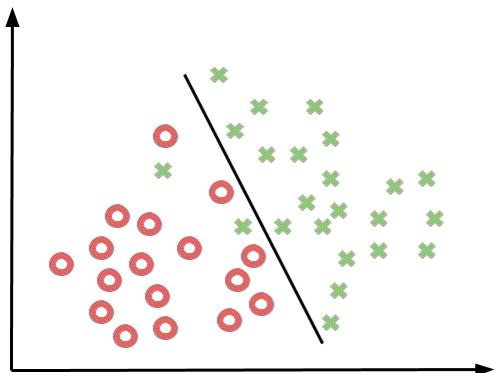
Unsupervised



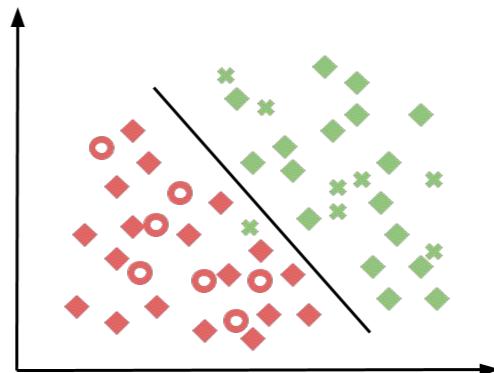
---

# Learning

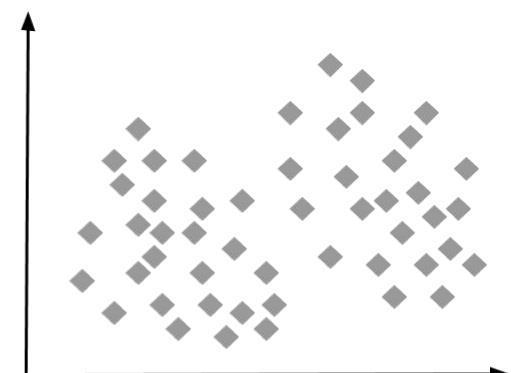
Supervised



Semi-supervised



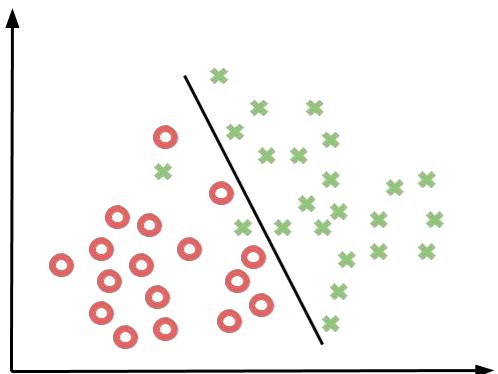
Unsupervised



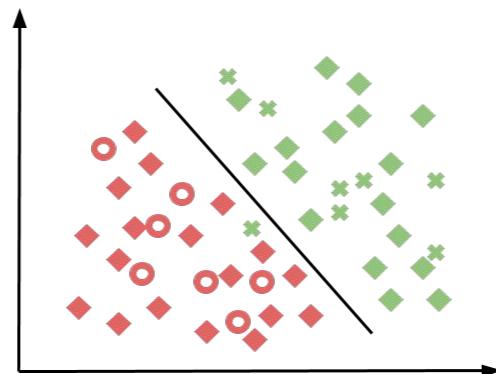
---

# Learning

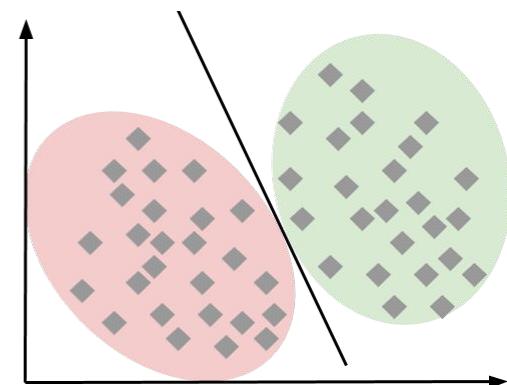
Supervised



Semi-supervised

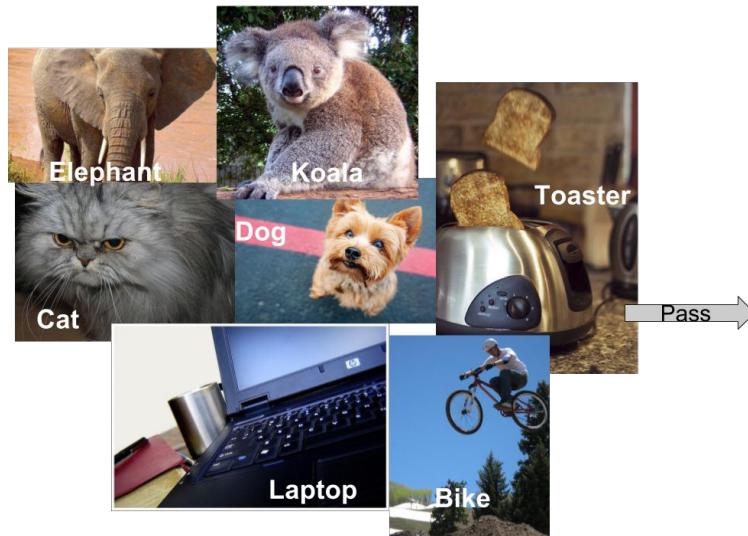


Unsupervised

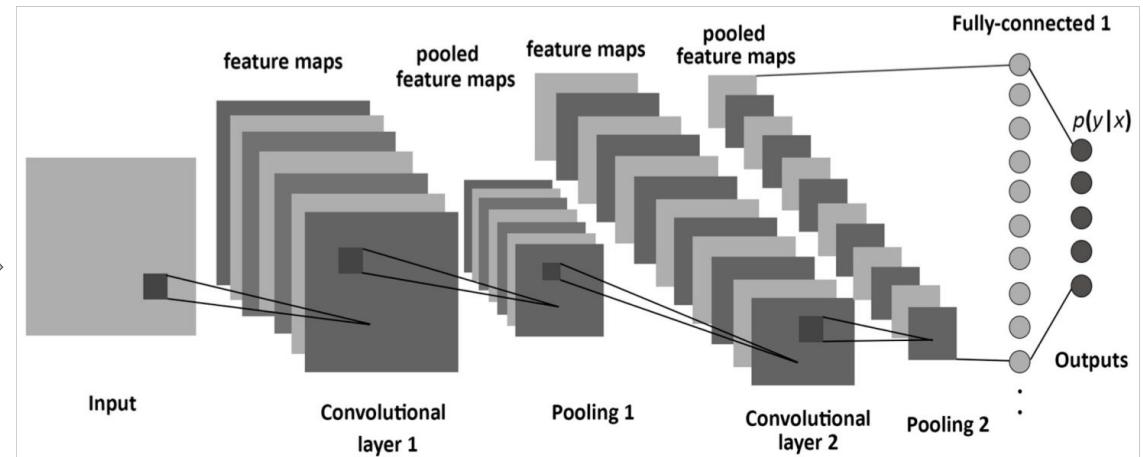


---

# Many-shots supervised learning



Pass

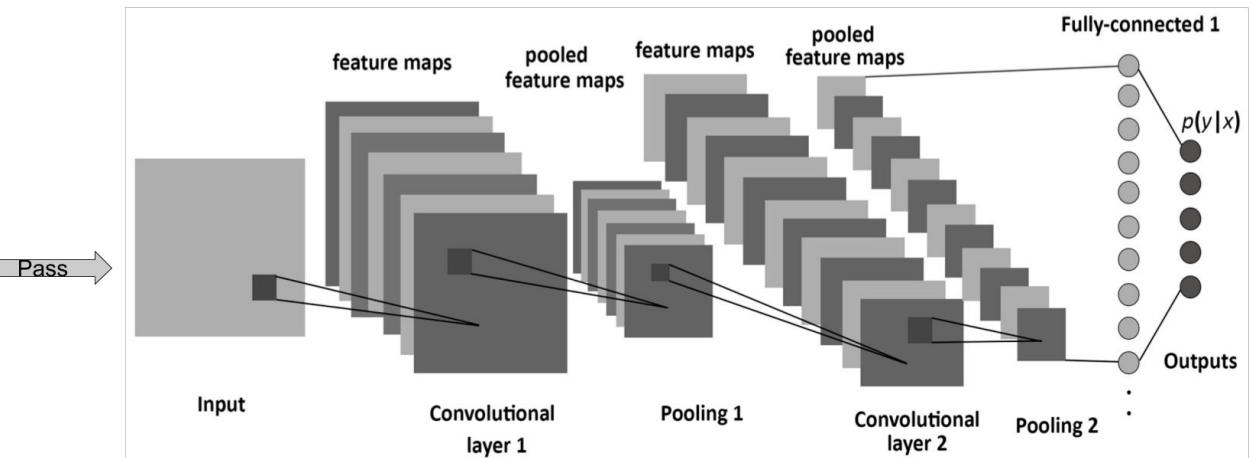


---

# Many-shots supervised learning



Huge amount of labeled data!



---

# Annotation Effort

IMAGENET



14M images, 21K categories

[Deng, et al CVPR2009]

# Annotation Effort

- OpenImages  
9M images, 6K categories
- COCO  
330K images
- MIT Places  
2.5M images



Complexity

# Task complexity

- Classification

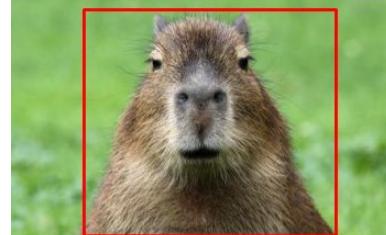
Single labeling is relatively easy task



Complexity ↓

# Task complexity

- Classification  
Single labeling is relatively easy
- Localization  
Requires precise bounding boxes



Complexity ↓

# Task complexity

- Classification

Single labeling is relatively easy

- Localization

Requires precise bounding boxes

- Captioning

Many other possibilities!

“Surprised capybara looking at the camera”

“Portrait photo of a capybara”



Complexity ↓

# Task complexity vs. annotation

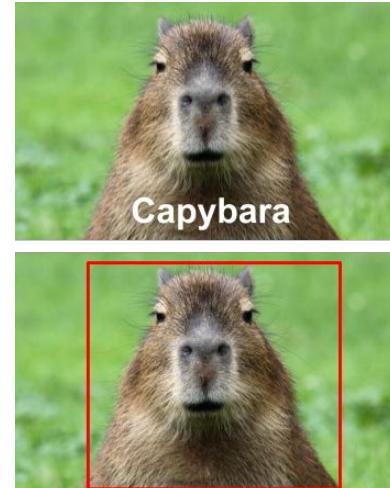
- Classification  
ImageNet 14M labeled images



Complexity ↓

# Task complexity vs. annotation

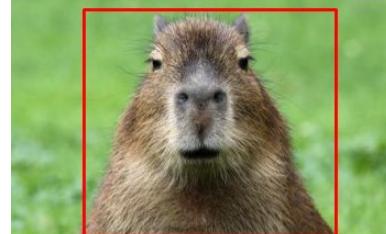
- Classification  
ImageNet 14M labeled images
- Localization  
openImages 3.6M bounding boxes



Complexity ↓

# Task complexity vs. annotation

- Classification  
ImageNet 14M labeled images
- Localization  
openImages 3.6M bounding boxes
- Captioning  
COCO 300K captioned images



Complexity ↓

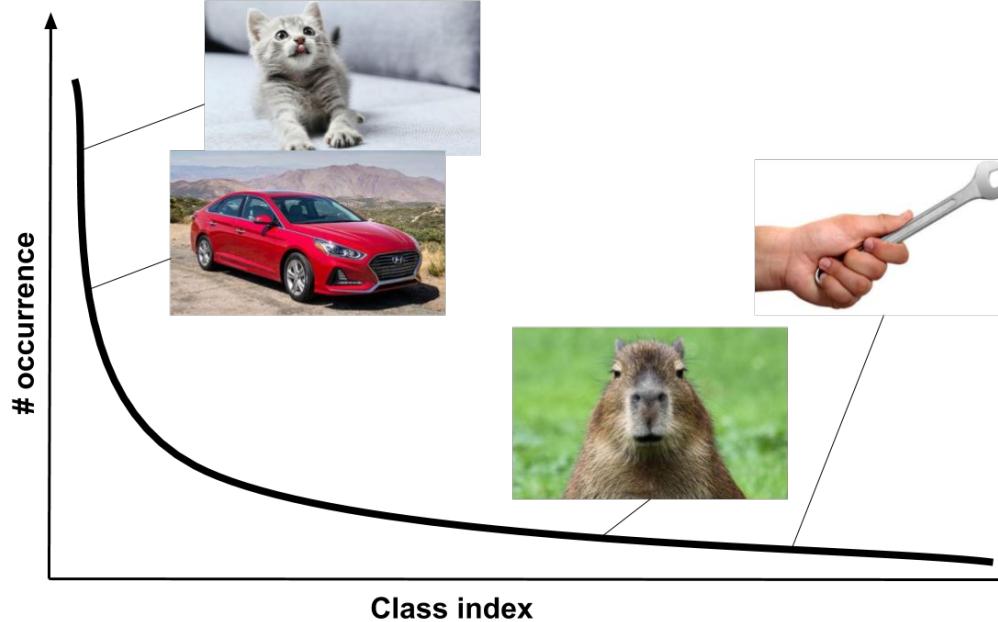
# Task complexity vs. annotation

- Classification  
ImageNet 14M labeled images
- Localization  
openImages 3.6M bounding boxes
- Captioning  
COCO 300K captioned images

**Less annotations for more complex tasks!**



# Tail distribution



<https://www.cars.com/>

<http://www.foxnews.com/lifestyle/2017/11/09/how-to-keep-cat-from-scratching-your-sofa-to-shreds.html>

<https://www.livescience.com/55223-capybara-facts.html>

<https://www.indiamart.com/proddetail/hand-wrench-13045857897.html>

---

# Fine-grained categories

Hard: subtle differences



Oxford Pet dataset

---

# Fine-grained categories

Hard: subtle differences



Cars dataset

---

# Fine-grained categories

Hard: subtle differences



---

# Fine-grained categories

Hard: subtle differences



60 classes of Caltech Birds dataset

---

# Fine-grained categories

Novice annotator



---

# Fine-grained categories

Bird expert → expensive





# Zero-shot Learning

---

# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises

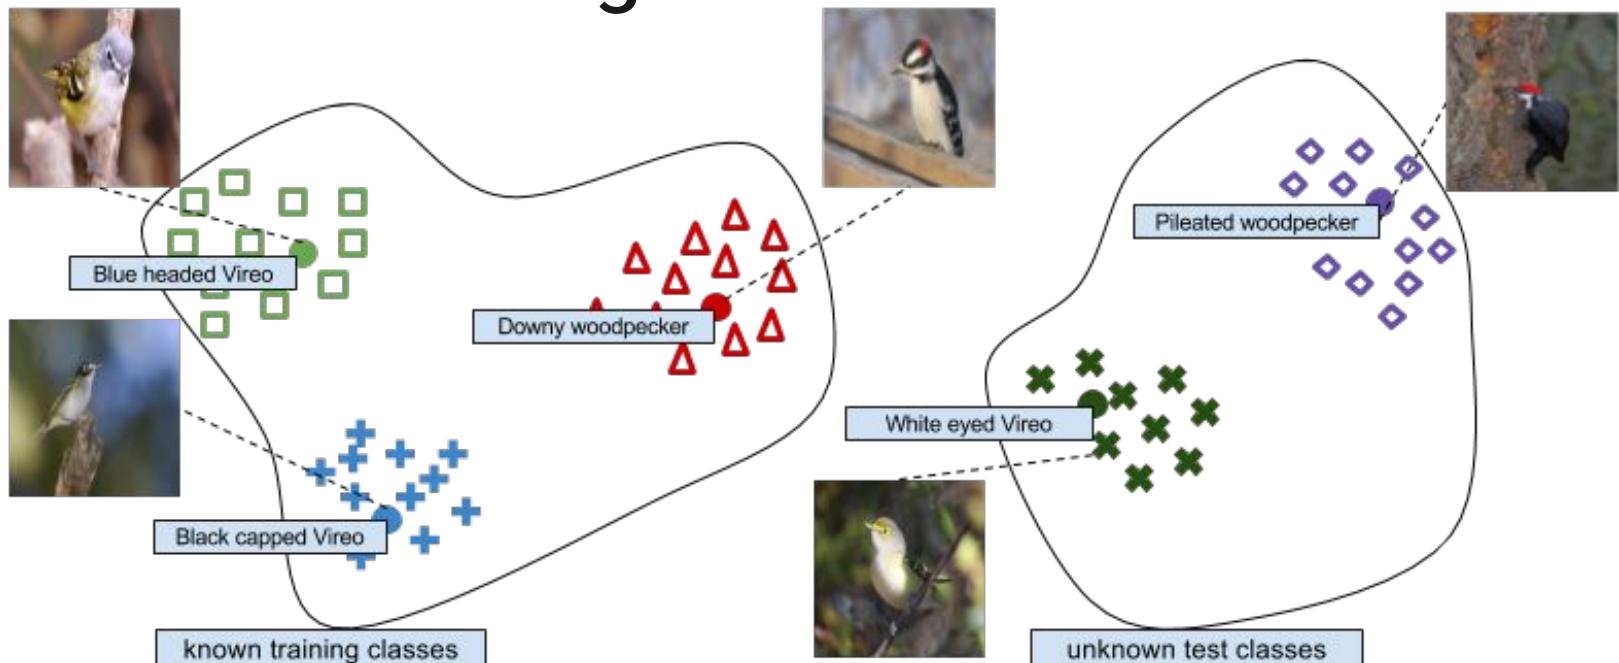
---

# Zero-shot learning

- Extreme case of scarce training data for some categories
- Disjoint sets of training and test classes
- New class examples appear only after training stage (at test time)

Animals classifier: No panda images at training time → panda images in testset

# Zero-shot learning



---



# **Zero-shot learning**

How?

---

# Zero-shot learning



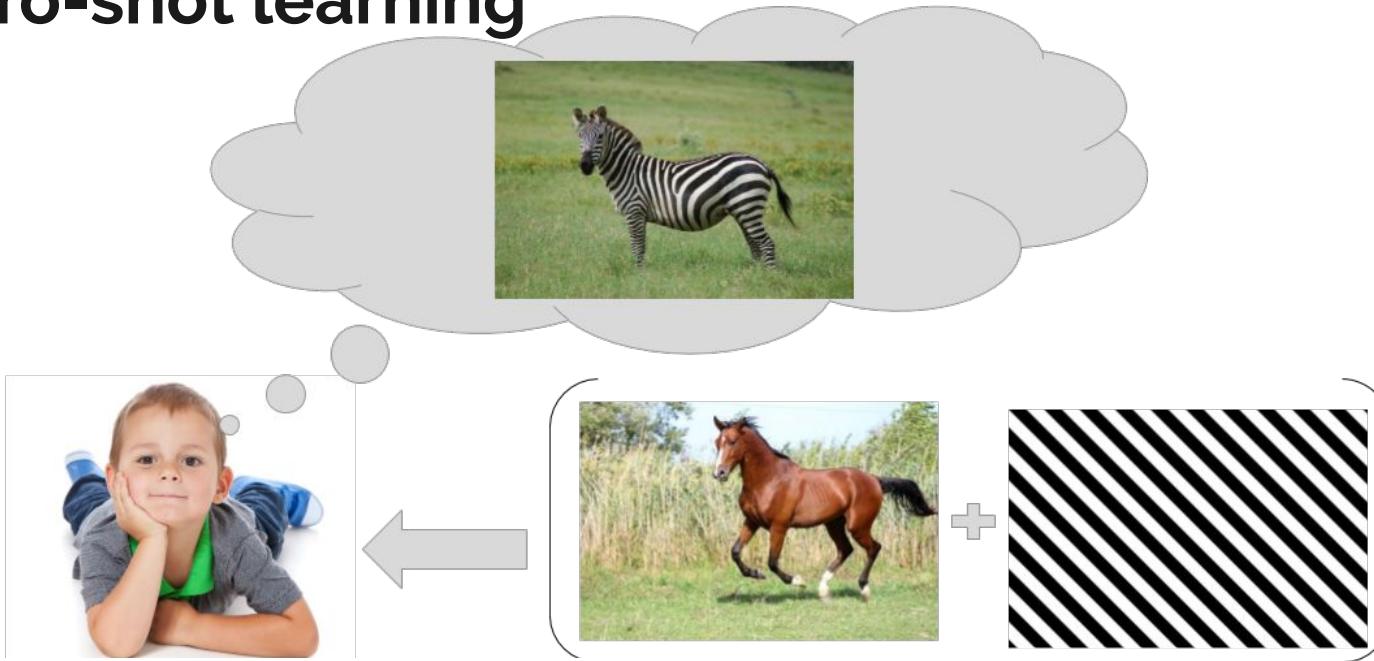
<http://www.youngparents.com.sg/development/10-steps-raising-self-confident-child/>

<http://www.thehorse.com/articles/33568/5-tips-for-packing-the-pounds-on-performance-horses>

<http://earlylearningtoys.org/stripes/>

---

# Zero-shot learning



<http://www.youngparents.com.sg/development/10-steps-raising-self-confident-child/>

<http://www.thehorse.com/articles/33568/5-tips-for-packing-the-pounds-on-performance-horses>

<http://earlylearningtoys.org/stripes/> <https://pt.wikipedia.org/wiki/Zebra-de-grant>

---

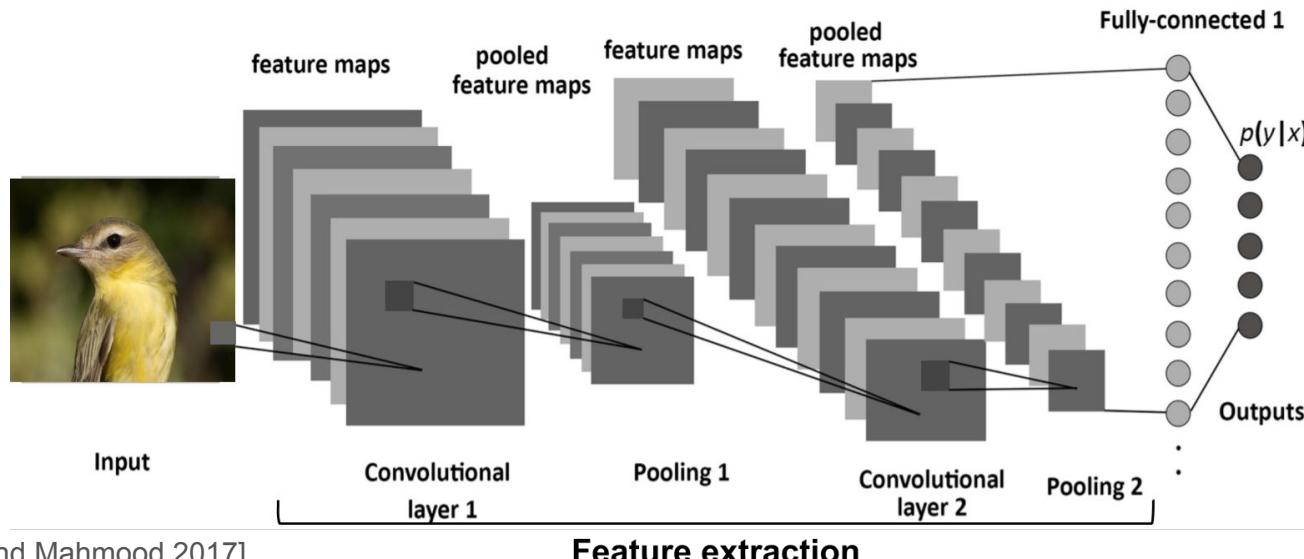


## Zero-shot learning

Knowledge transfer & Side information

# Knowledge transfer

Use bottleneck image features of pre-trained model





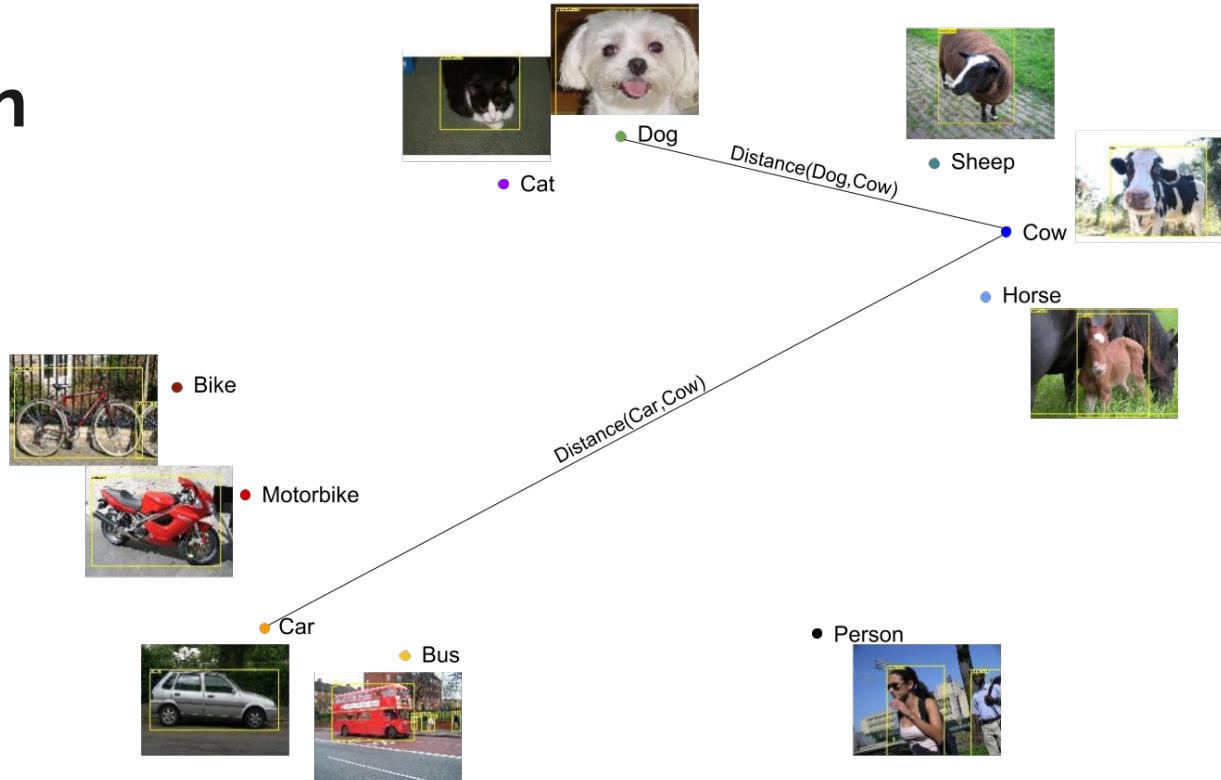
# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises

# Side information

Use side information to get a better representation for classes (categories).

Labels → Vector representations



---

# Side information - Attributes embedding



**Polar bear**

black: no  
white: yes  
brown: no  
stripes: no  
water: yes  
eats fish: yes

[0 1 0 0 1 1]



**Otter**

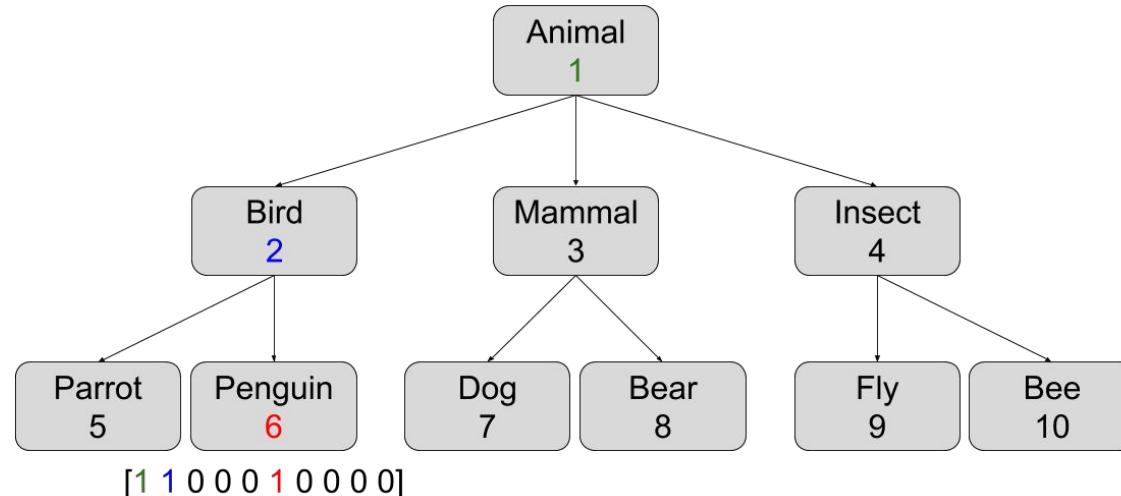
black: yes  
white: no  
brown: yes  
stripes: no  
water: yes  
eats fish: yes

[1 0 1 0 1 1]

---

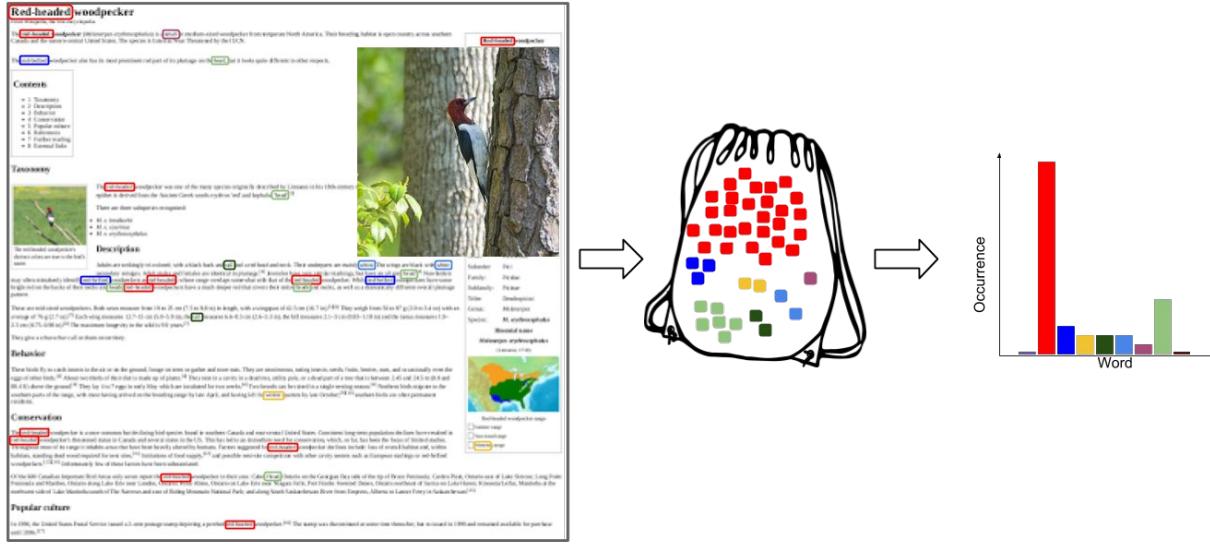
# Side information - Hierarchical embedding

HLE Hierarchical Label Embedding extracted from Wordnet



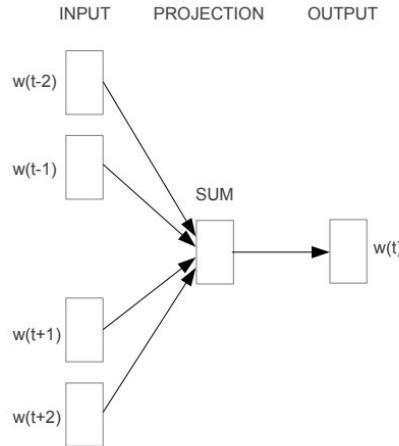
# Side information - Text embedding

## Sparse representations: BoW Bag of Words from Wikipedia articles

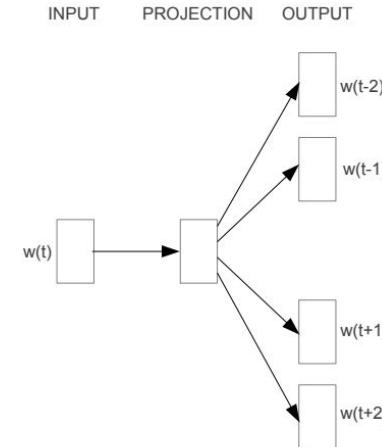


# Side information - Text embedding

Dense representations: Word2vec



CBOW



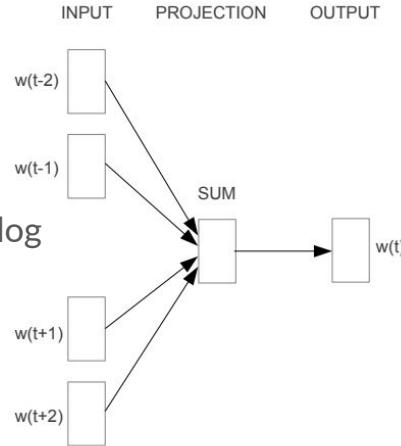
Skip-gram

[Mikolov, et al. NIPS2013]

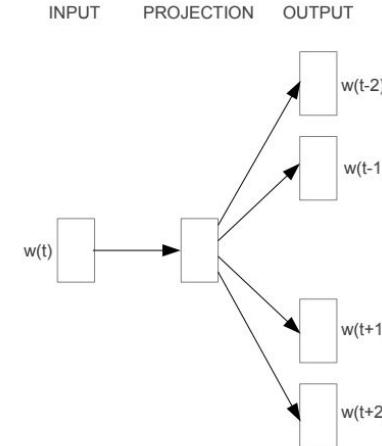
# Side information - Text embedding

Dense representations: Word2vec

The fox \_\_?\_\_ over the lazy dog



CBOW



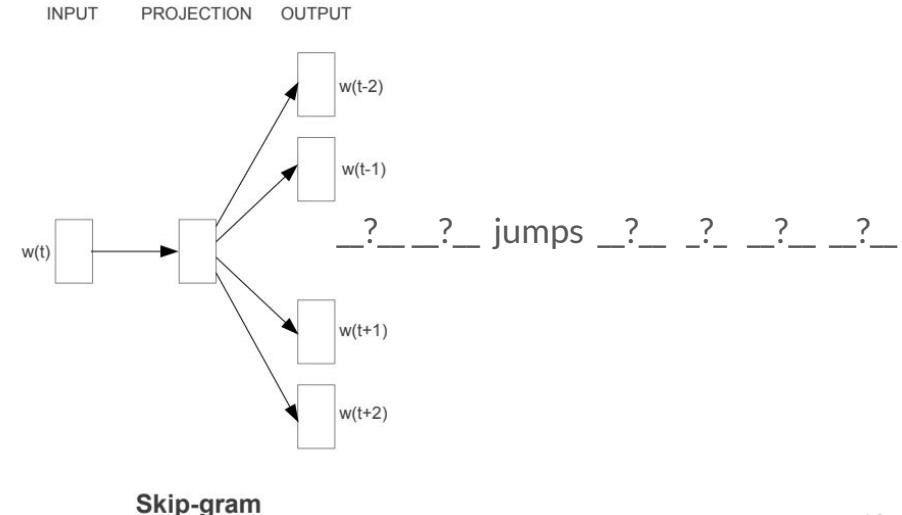
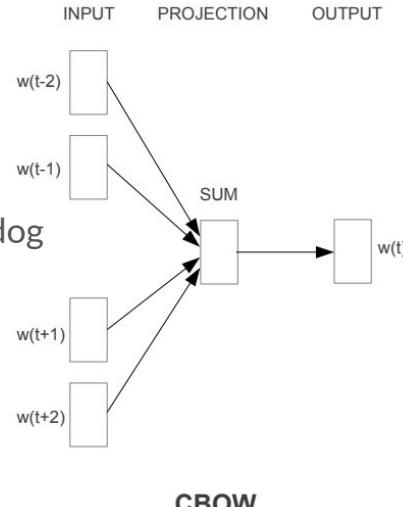
Skip-gram

[Mikolov, et al. NIPS2013]

# Side information - Text embedding

Dense representations: Word2vec

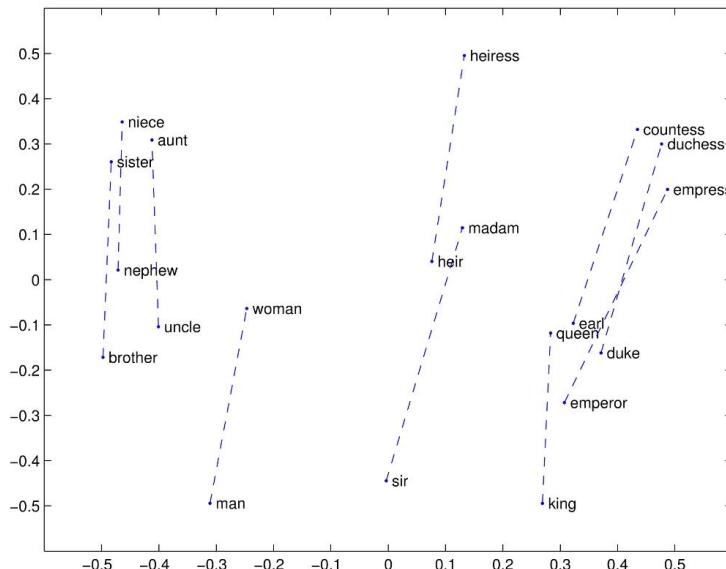
The fox \_?\_ over the lazy dog



---

# Side information - Text embedding

Dense representations: Glove Global vectors for word representation



---

## Side information - Visual descriptions



The bird has a white underbelly, black feathers in the wings, a large wingspan, and a white beak.



This flower has a central white blossom surrounded by large pointed red petals which are veined and leaflike.



This bird has distinctive-looking brown and white stripes all over its body, and its brown tail sticks up.

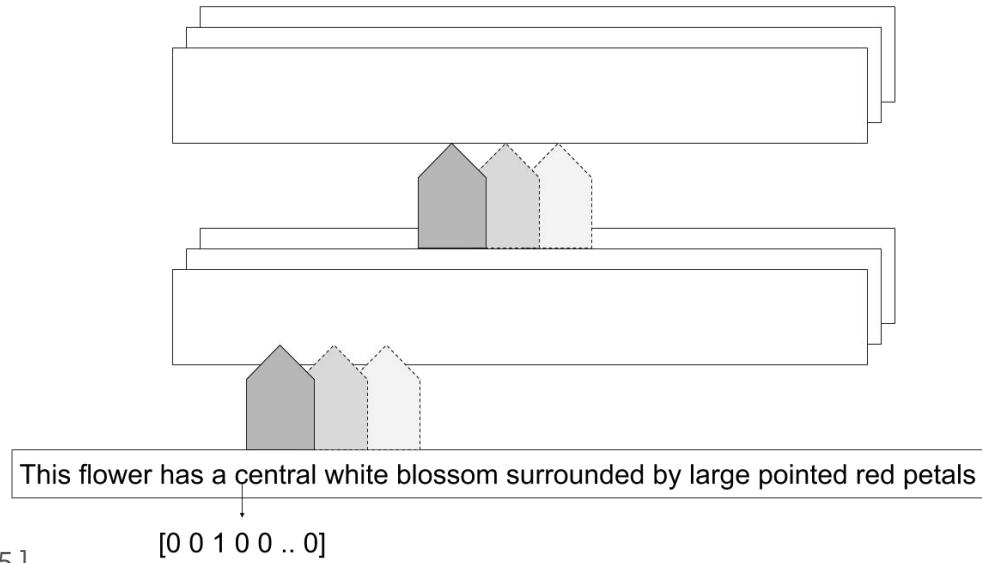


Light purple petals with orange and black middle green leaves

---

# Side information - Visual descriptions

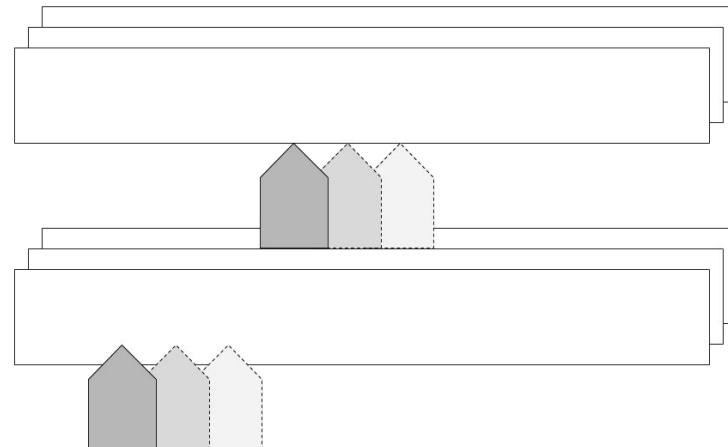
- Character level CNN



---

# Side information - Visual descriptions

- Word level CNN



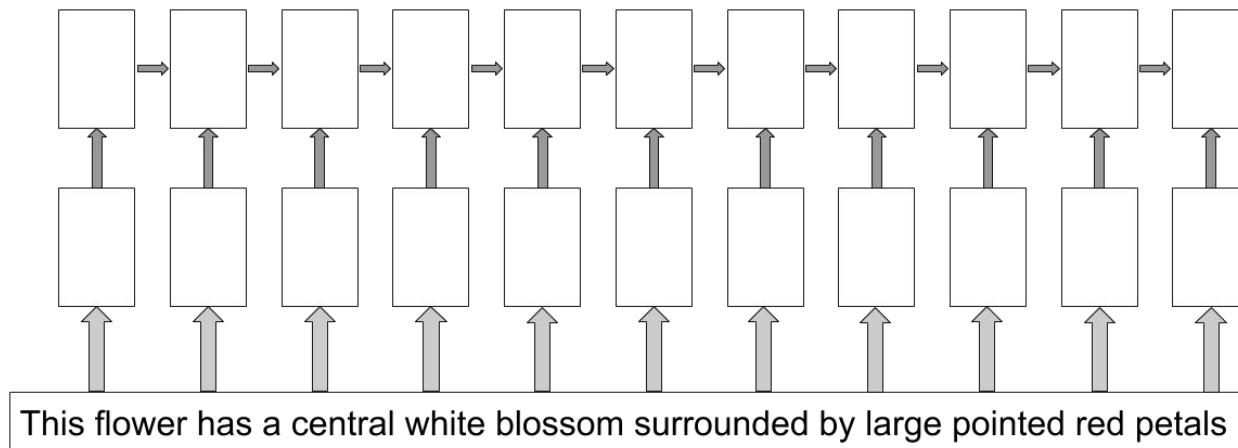
This flower has a central white blossom surrounded by large pointed red petals

[0 0 .. 1 0 0]

---

# Side information - Visual descriptions

- LSTM



---

## Side information - Gaze embedding

- Discrimination of objects by novice
- Data collection is fast
- Implicit annotation,  
you don't need to name the object



---

# Side information - Gaze embedding

Experiment

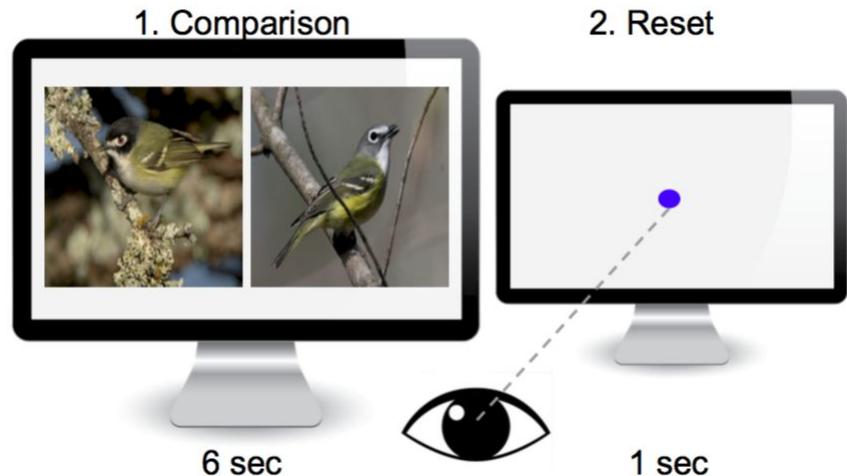
1. Comparison



---

# Side information - Gaze embedding

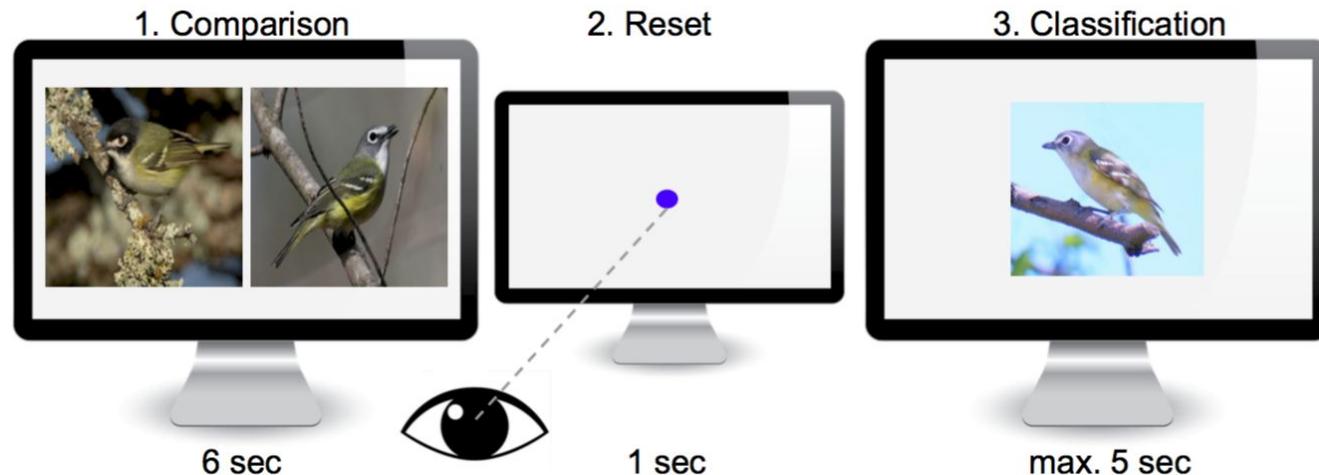
Experiment



---

# Side information - Gaze embedding

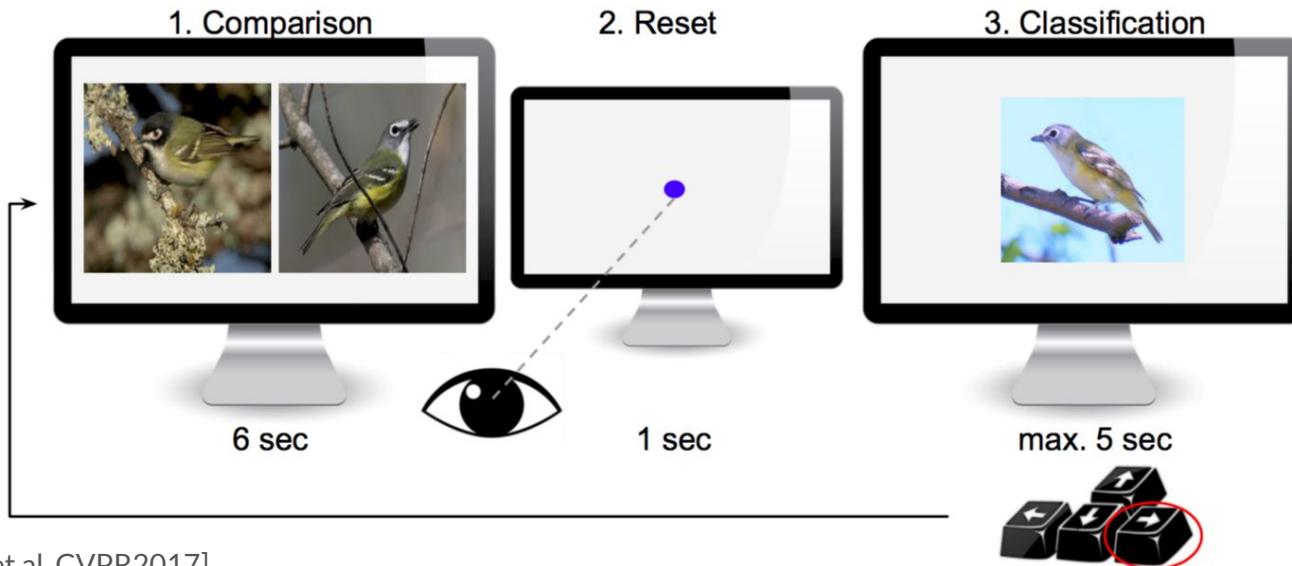
Experiment



---

# Side information - Gaze embedding

Experiment

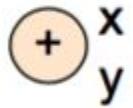


---

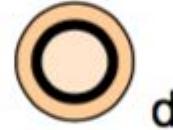
# Side information - Gaze embedding

Features

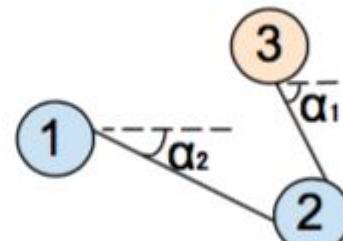
Location



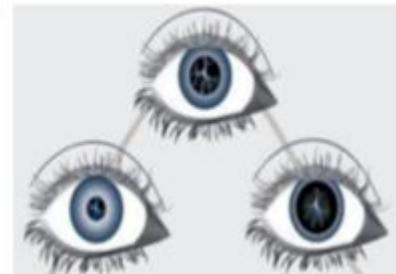
Duration



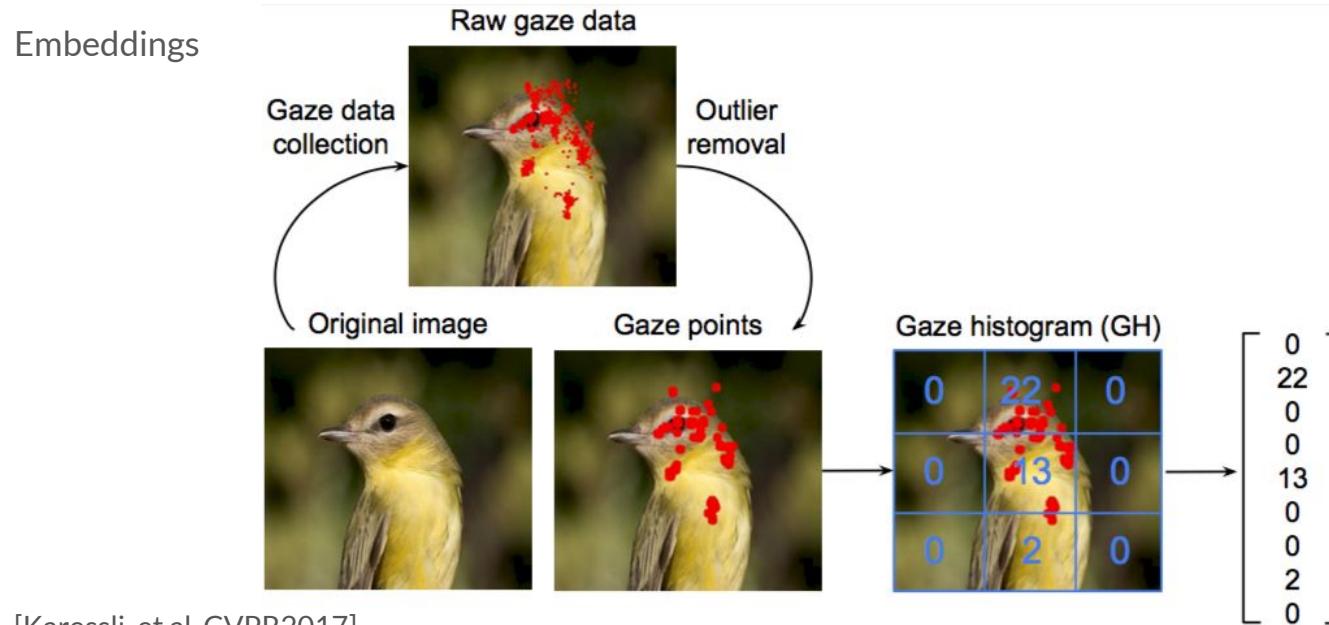
Sequence



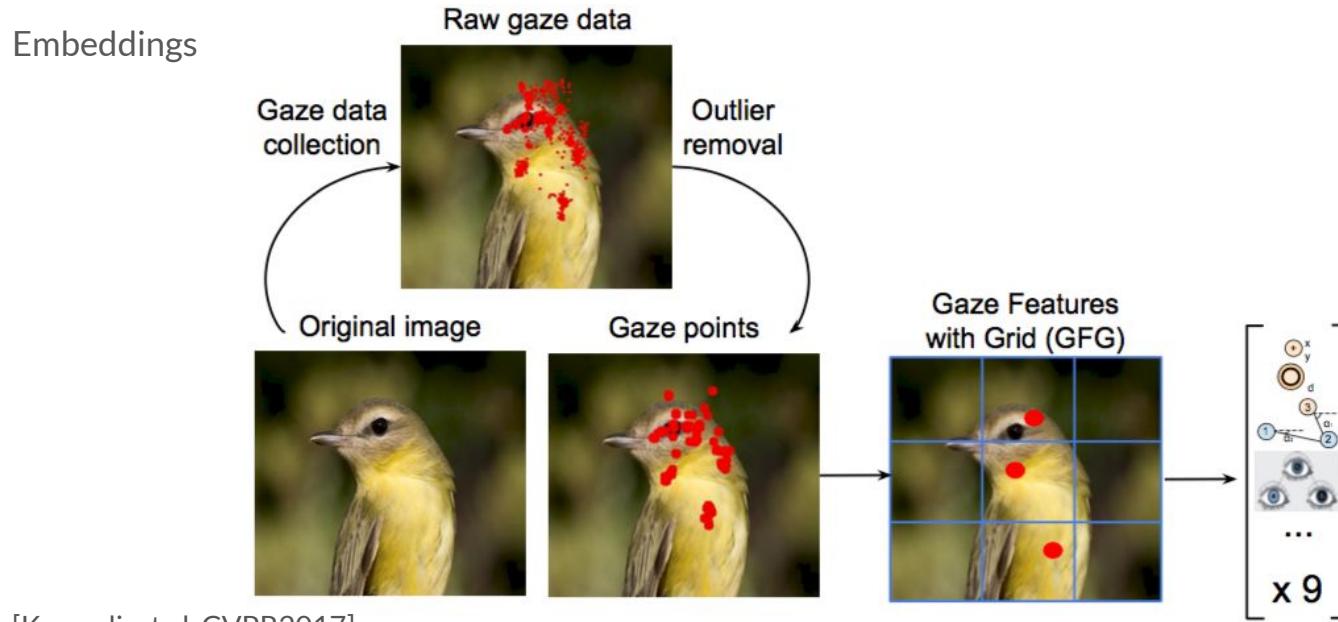
Pupil Diameter



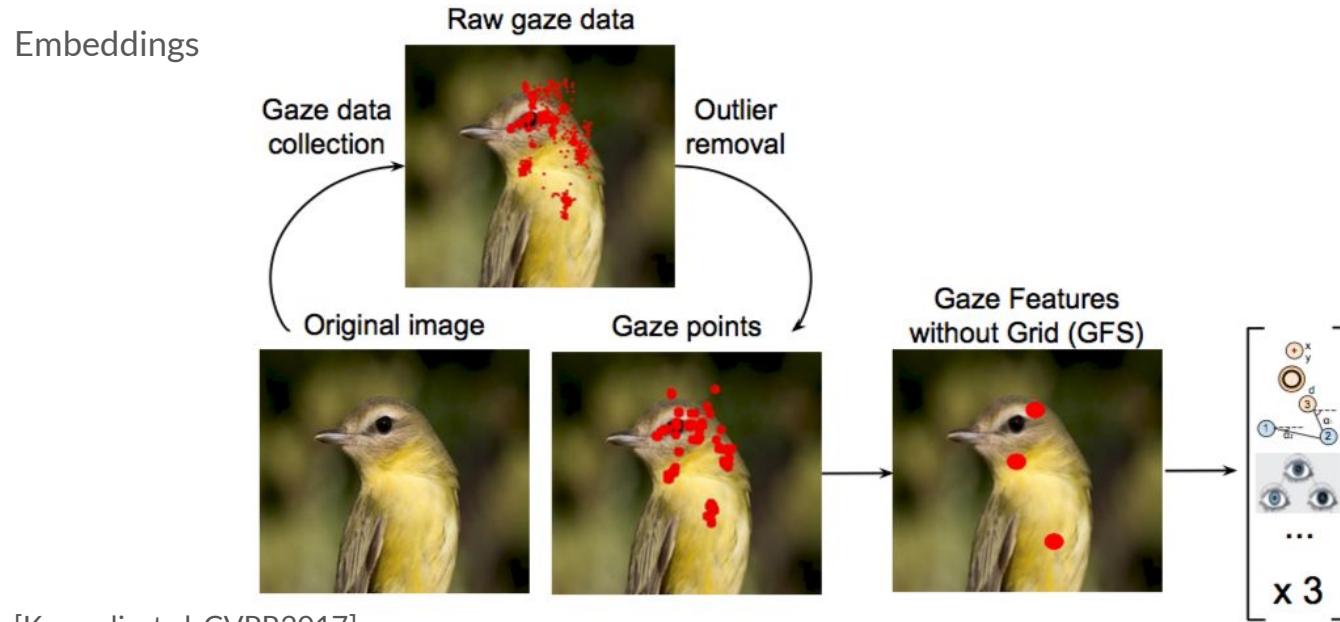
# Side information - Gaze embedding



# Side information - Gaze embedding

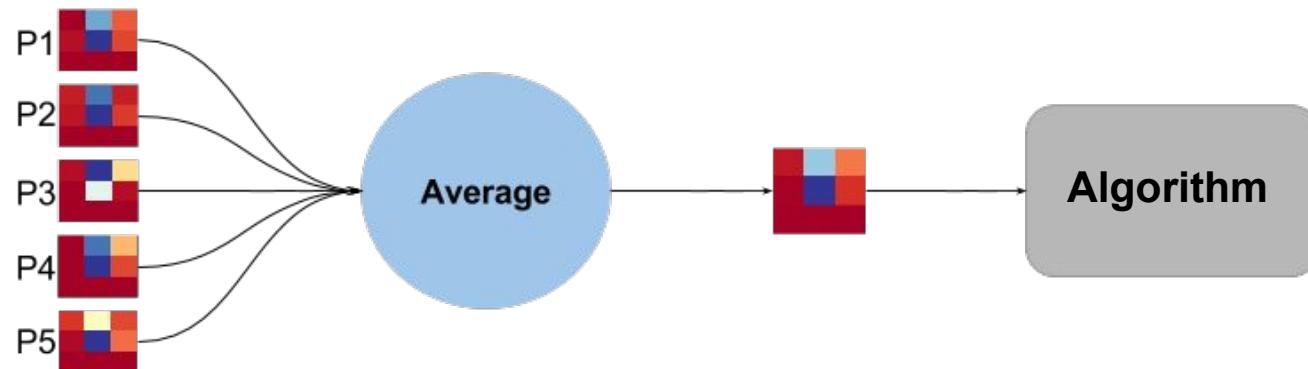


# Side information - Gaze embedding



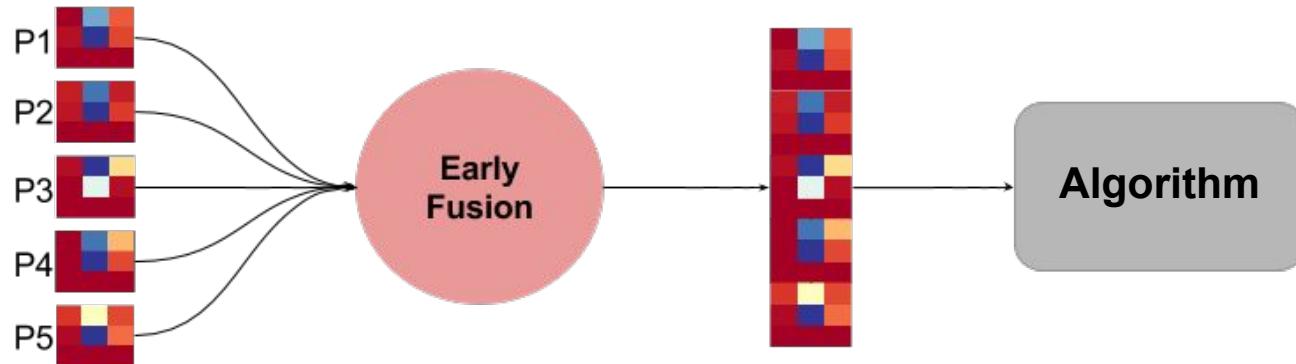
---

## Side information - Gaze embedding



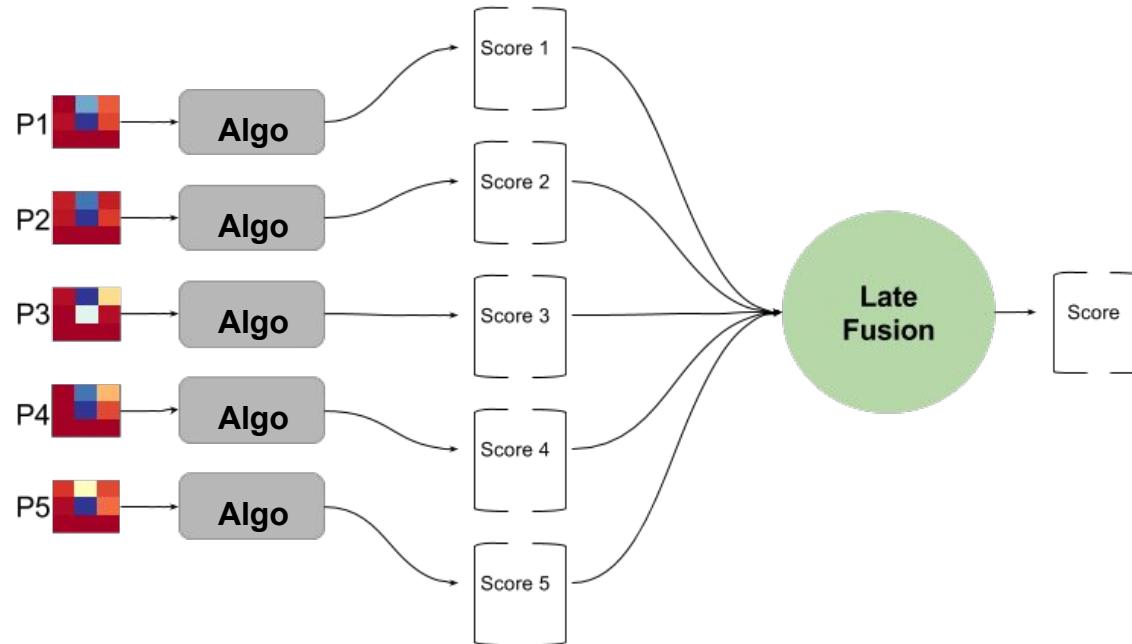
---

# Side information - Gaze embedding



---

## Side information - Gaze embedding



---

# Side information summary

- Supervised
  - Expert
    - Class attributes
  - Novice
    - Detailed visual descriptions (deep representations of visual descriptions)
    - Human gaze
- Unsupervised
  - Hierarchical similarity
  - Text embeddings

---

# CUB birds

- 11,788 images
- 200 bird species, 150 train+val set and 50 test classes



[Welinder, et al. 2010]

---

# Results

Source	Side information	Accuracy
Text	Hierarchical	20.6
	Bag of Words	22.1
	Word2vec	28.4
	Glove	24.2
Expert annotator	Attributes	50.1
Novice annotator	Detailed visual descriptions	<b>56.8</b>

---

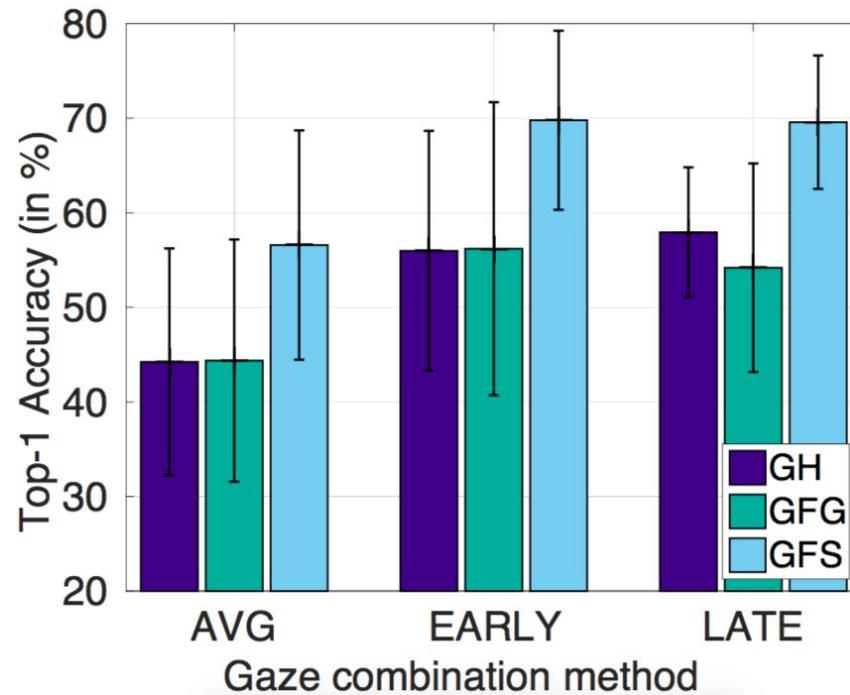
# Mini-CUB birds

- 464 images
- 14 bird species, 11 train+val set and 3 test classes



---

# Results



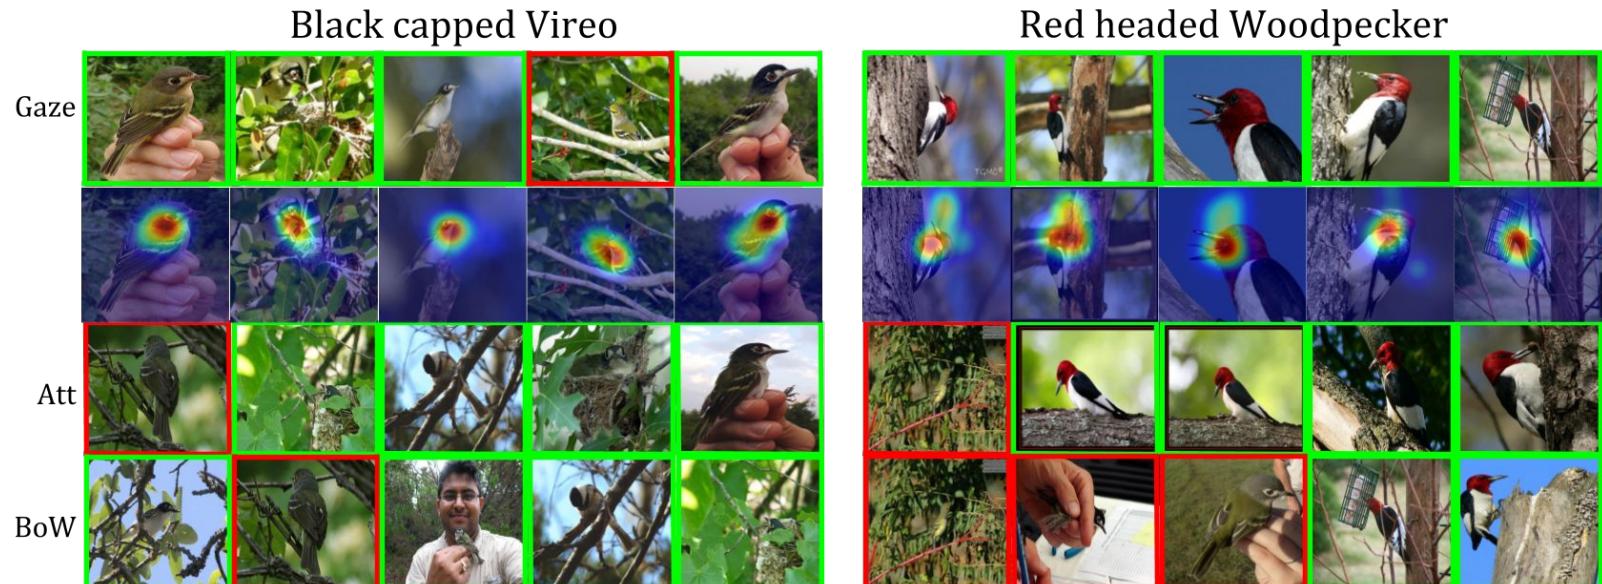


# Results

Method	Accuracy
Bag-of-Words from Wiki	55.2
Human annotated attributes	72.9
Gaze embeddings	73.9
Attributes + Gaze	<b>78.2</b>

---

# Results



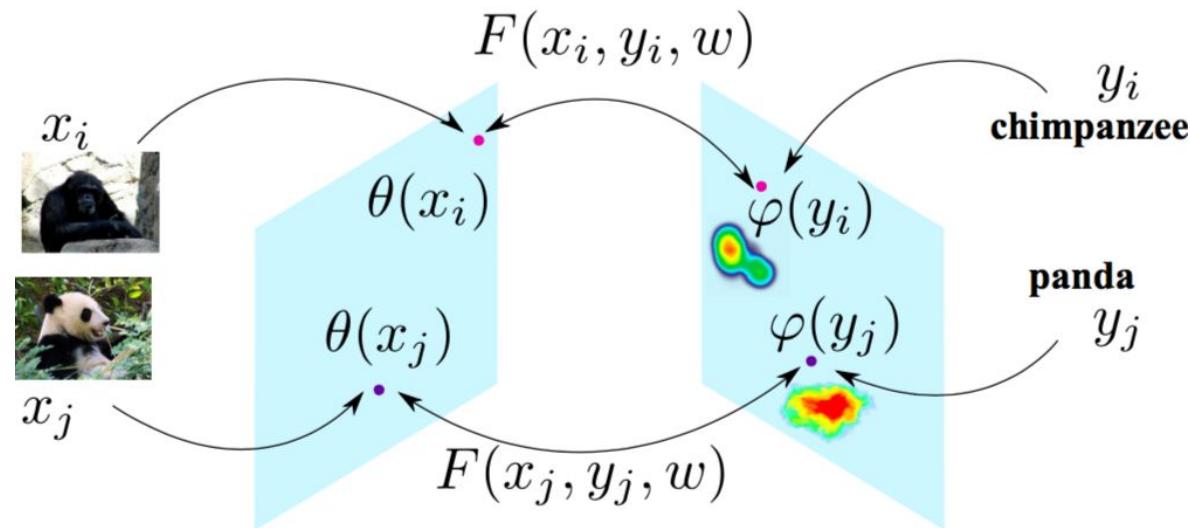


# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises

---

# Zero-shot models



---

# Task formulation

Training set  $S = \{(x_n, y_n), n = 1..N\}, y_n \in Y^{train}$

We want to learn a function  $f : X \rightarrow Y$  by minimize the empirical risk:

$$\frac{1}{N} \sum_{n=1}^N L(y_n, f(x_n; W)) + \Omega(W)$$

- Loss function
- Regularization term

---

# Task formulation

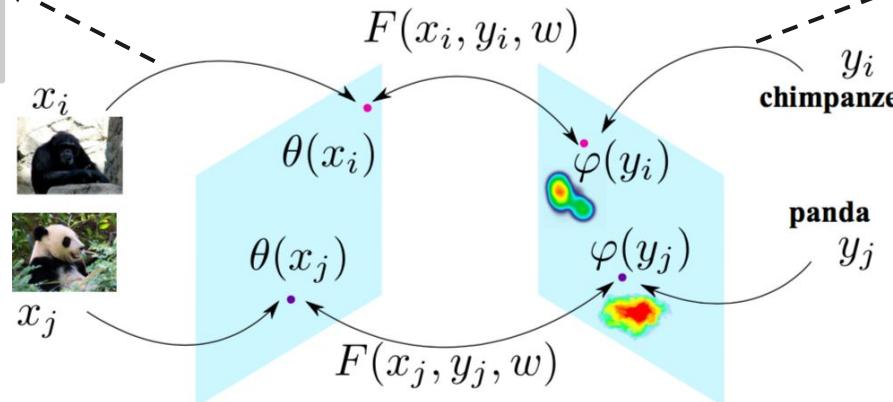
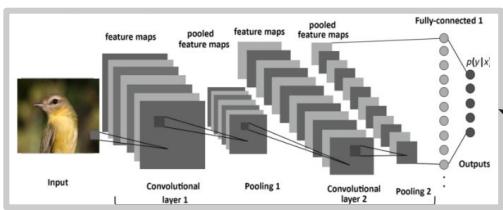
Training set  $S = \{(x_n, y_n), n = 1..N\}, y_n \in Y^{train}$

Maximizing the compatibility (score):

$$f(x_n; W) = \operatorname{argmax}_{y \in Y} F(x, y; W)$$

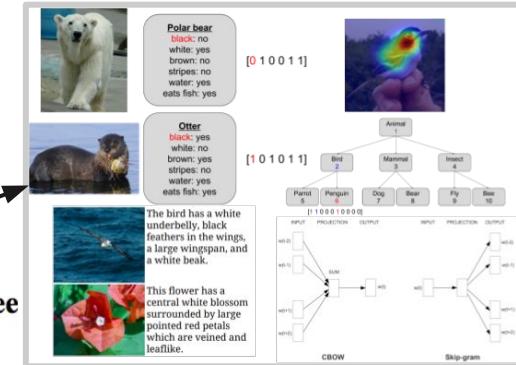
At test time we predict  $y \in Y^{test}$  that gives the highest compatibility  $Y^{test} \subset Y$

# Linear Compatibility



[Akata, et al. CVPR2015]

$$F(x, y; W) = \theta(x)^T W \phi(y)$$

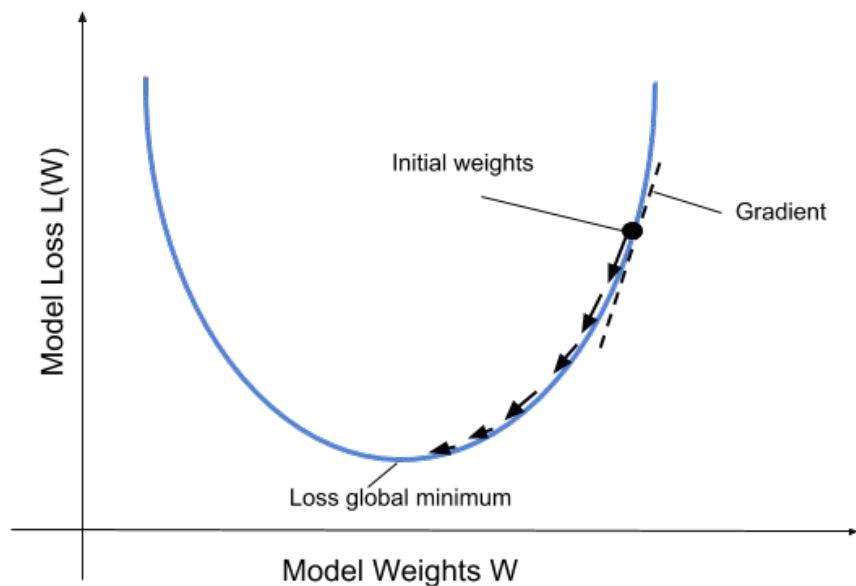


# Recap - Gradient Descent

Minimize the loss function by updating the weights with respect to the gradient.

**Gradient descent**  
updates weights only after each epoch

**Stochastic gradient descent**  
updates weights on each sample



---

# Linear Compatibility - DEVISE

Deep visual semantic embedding, pairwise ranking

$$\sum_{y \in \mathcal{Y}^{tr}} [\Delta(y_n, y) + F(x_n, y; W) - F(x_n, y_n; W)]_+$$

$$\Delta(y_n, y) = \begin{cases} 0, & \text{if } y_n = y. \\ 1, & \text{otherwise.} \end{cases}$$

---

## Linear Compatibility - ALE

Attribute label embedding, weighted pairwise ranking

$$\sum_{y \in y^{tr}} l_k [\Delta(y_n, y) + F(x_n, y; W) - F(x_n, y_n; W)]_+$$

$$l_k = \sum_{i=1}^k \alpha_i \quad \text{where} \quad \alpha_i = 1/i$$

---

# Linear Compatibility - SJE

Structured joint embedding, multiclass objective

$$[\max_{y \in \mathcal{Y}^{tr}} (\Delta(y_n, y) + F(x_n, y; W)) - F(x_n, y_n; W)]_+$$



# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises



# Zero-shot learning exercise



# **Recap Exercise - Linear Regression & Gradient Descent**

---

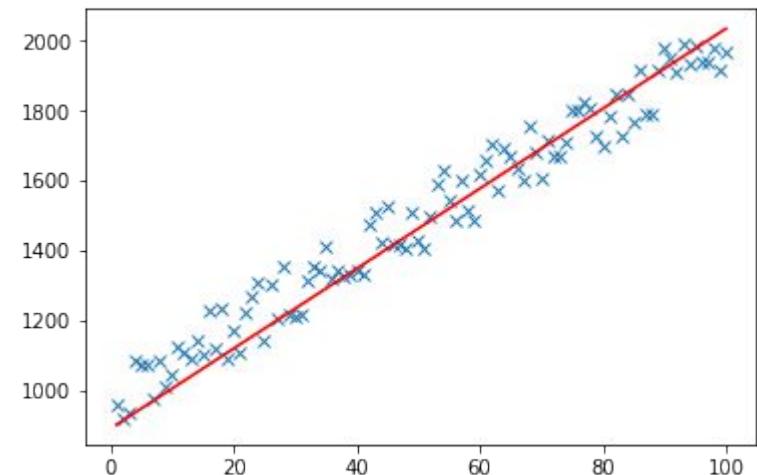
# Linear regression - implementation

- Clone [repo](#)
- Generate data for linear regression
- Implement optimisation GD and SGD

Linear regression

$$y = w_0 x_0 + w_1 x_1 ;$$

$$x_0 = 1$$

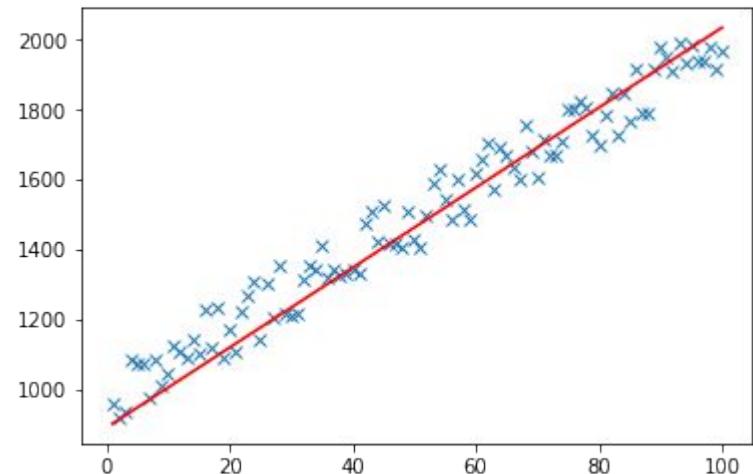


# Linear regression with GD - implementation

$$\text{Loss} = \frac{1}{2N} (P - Y)^2$$

**foreach epoch**

1. Calculate predictions  $P = X * W$
2. Calculate error  $= (P - Y)^2$
3. Calculate gradient  $= X.T * (P - Y) / N$
4. Update weights  $W = W - lr * \text{gradient}$

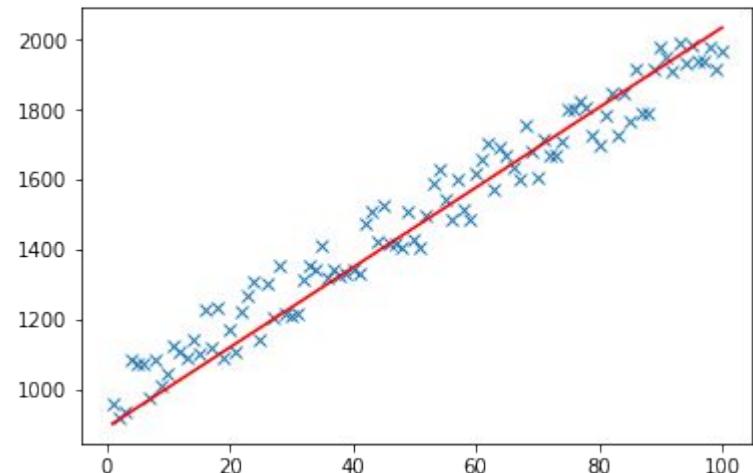


# Linear regression with SGD - implementation

$$\text{Loss} = \frac{1}{2N} (P-Y)^2$$

**foreach** epoch

1. Shuffle training samples
2. **foreach** sample  $x_i, y_i$ 
  1. Calculate sample prediction  $p_i = x_i * W$
  2. Calculate error $_i = (p_i - y_i)^2$
  3. Calculate gradient $_i = x_i * (p_i - y_i)$
  4. Update weights  $W = W - lr * \text{gradient}_i$





# **Structured Joint Embedding exercise**

---

# AWA dataset

Animal with attributes dataset

- 30K images
- 50 classes
- 85 attribute
- Standard split 40 train+val 10 test

---

# SJE implementation

- Select dataset (AWA)
- Download [data](#) (image features & corresponding labels) & use class embeddings in [repo](#)
- Implement zero-shot algorithm SJE
- Evaluate Top-1 accuracy on test set

# SJE implementation

Initialize W (DxE)

**foreach** epoch

    Shuffle training samples

**foreach** sample ( $x_i$ ,  $y_{true}$ )

        score<sub>true</sub> =  $\theta(x_i) * W^T * \phi(y_{true})$

        loss<sub>max</sub> = -1,  $y_{max}$  = -1

**foreach** training class  $y_{tr}$

          score =  $\theta(x_i) * W^T * \phi(y_{tr})$

          loss =  $\Delta(y_{tr}, y_{true}) + score - score_{true}$

**if** loss > loss<sub>max</sub>

              update loss<sub>max</sub> and  $y_{max}$

**if**  $y_{max} \neq y_{true}$

$W = W - lr * (\theta(x_i) \otimes (\phi(y_{max}) - \phi(y_{true})) )$

D: x dim, E: y dim

find max score

violation  
update weights



# More Zero-shot models

---

## Linear Compatibility - ESZSL

Embarrassingly simple zero-shot learning, adds regularization terms

$$\gamma \| W\phi(y) \|^2 + \lambda \| \theta(x)^T W \|^2 + \beta \| W \|^2$$

$\gamma, \lambda, \beta$  regularization parameters

---

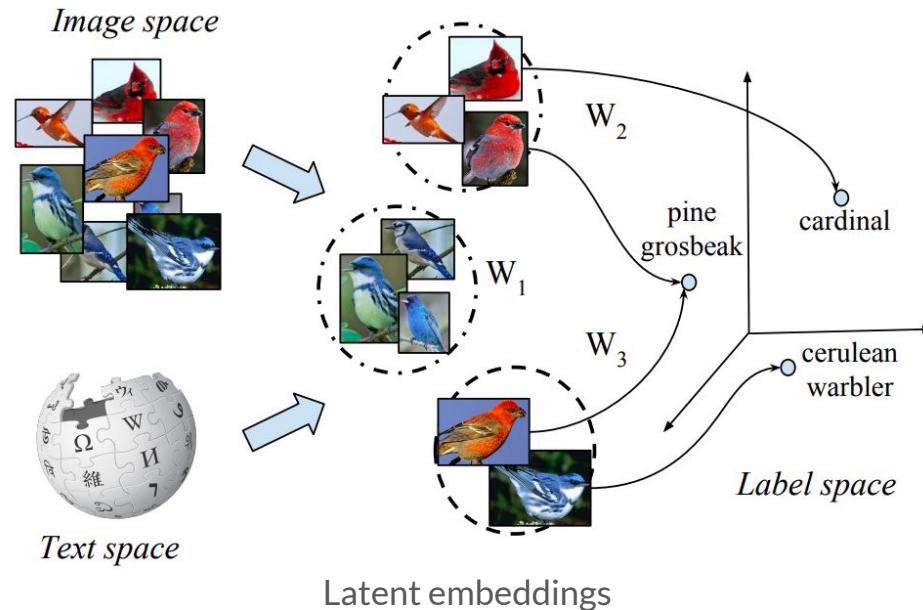
# Linear Compatibility - SAE

Semantic autoencoder, linear objective

$$\min_W \| \theta(x) - W^T \phi(y) \|^2 + \lambda \| W\theta(x) - \phi(y) \|^2$$

- Autoencoder learns projection from image features to label embedding
- The autoencoder must reconstruct original image features

# Nonlinear Compatibility - LATEM



---

# Nonlinear Compatibility - LATEM

Latent embeddings, piecewise linear compatibility

$$F(x, y; W_i) = \max_{1 \leq i \leq K} \theta(x)^T W_i \phi(y)$$

- Support non-linearity
- Different  $W$  encodes different characteristics

---

# Nonlinear Compatibility - CMT

Cross modal transfer, nonlinear compatibility using two layers NN

$$\sum_{y \in \mathcal{Y}^{tr}} \sum_{x \in \mathcal{X}_y} \|\phi(y) - W_1 \tanh(W_2 \cdot \theta(x))\|$$

# Intermediate Classifier - DAP

Direct attribute prediction

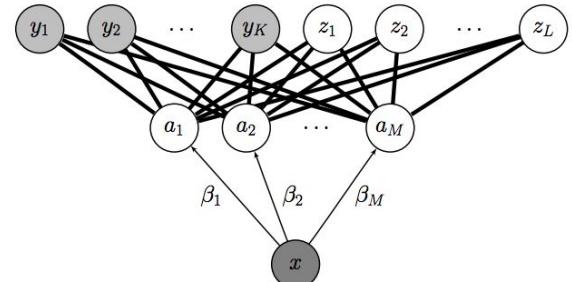
1. Learning probabilistic classifiers for each attribute
2. Combine scores

$$f(x) = \operatorname{argmax}_c \prod_{m=1}^M \frac{p(a_m^c | x)}{p(a_m^c)}$$



**Polar bear**

black: no  
white: yes  
brown: no  
stripes: no  
water: yes  
eats fish: yes



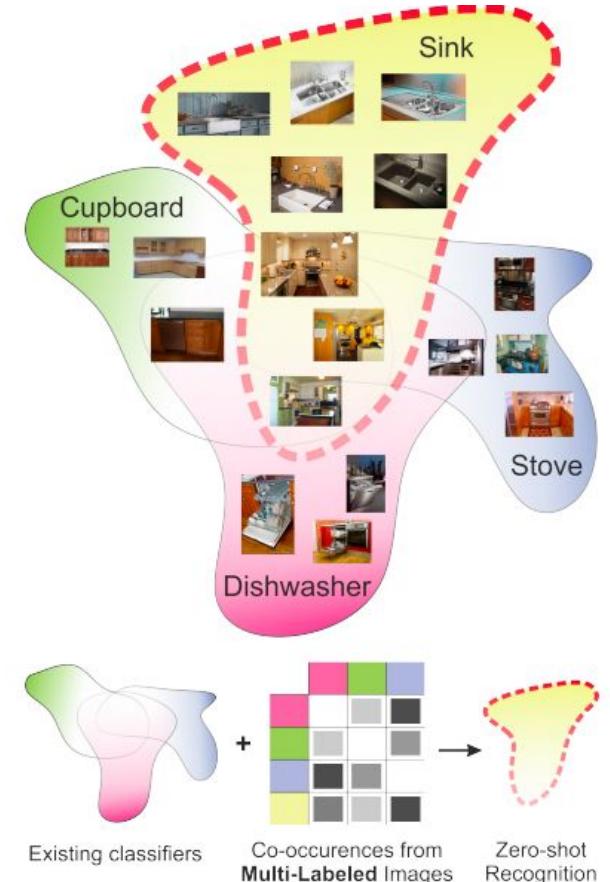
# Intermediate Classifier - COSTA

Co-occurrence statistics of visual concepts  
between seen and unseen classes

$$\hat{\mathbf{w}}_l = \sum_k \mathbf{w}_k s_{lk}$$

k classifier weights (svm)

Similarity between seen k and unseen l





# Models summary

## 1. Linear compatibility

DEVISE [Frome, et al. NIPS2013], ALE [Akata, et al. CVPR2015], SJE [Akata, et al. CVPR2015],  
ESZSL [Romera-Paredes and Torr ICML2015], SAE [Kodirov, et al. CVPR2017]

## 2. Nonlinear compatibility

LATEM [Xian, et al. CVPR2016], CMT [Socher, et al. NIPS2013]

## 3. Intermediate classifier

DAP [Lampert, et al. CVPR2009], COSTA [Mensink, et al. CVPR2014]



# AWA dataset

Animal with attributes dataset

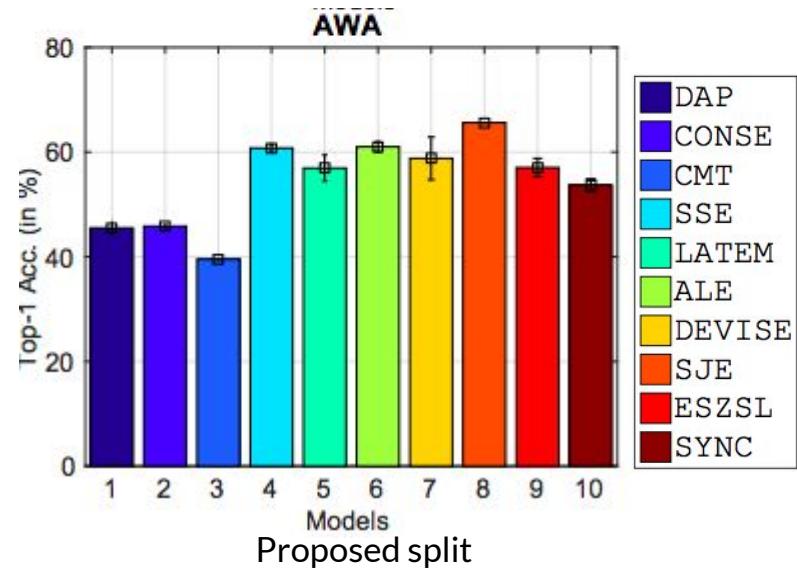
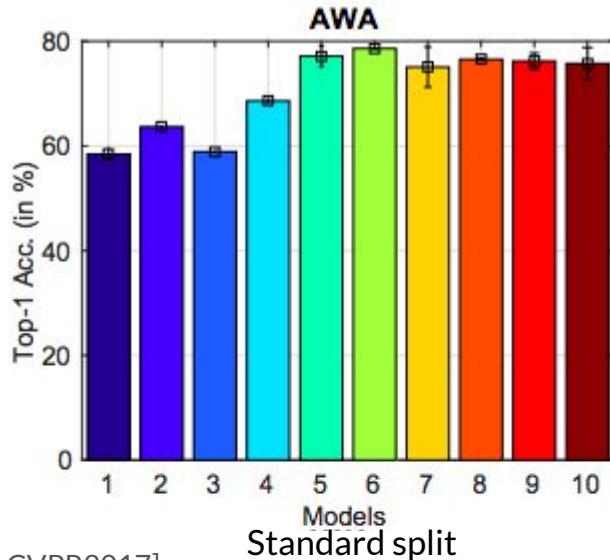
- 30K images
- 50 classes
- 85 attribute
- Standard split 40 train+val 10 test
- Suggested split

Insures that none of the test classes is used to train the image features base model

---

# Results

Evaluation on Animals with Attributes AWA





# Q & A