



# Learning with Small Samples

## Including zero-shot learning

Nour Karessli  
DSR 2018



# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises



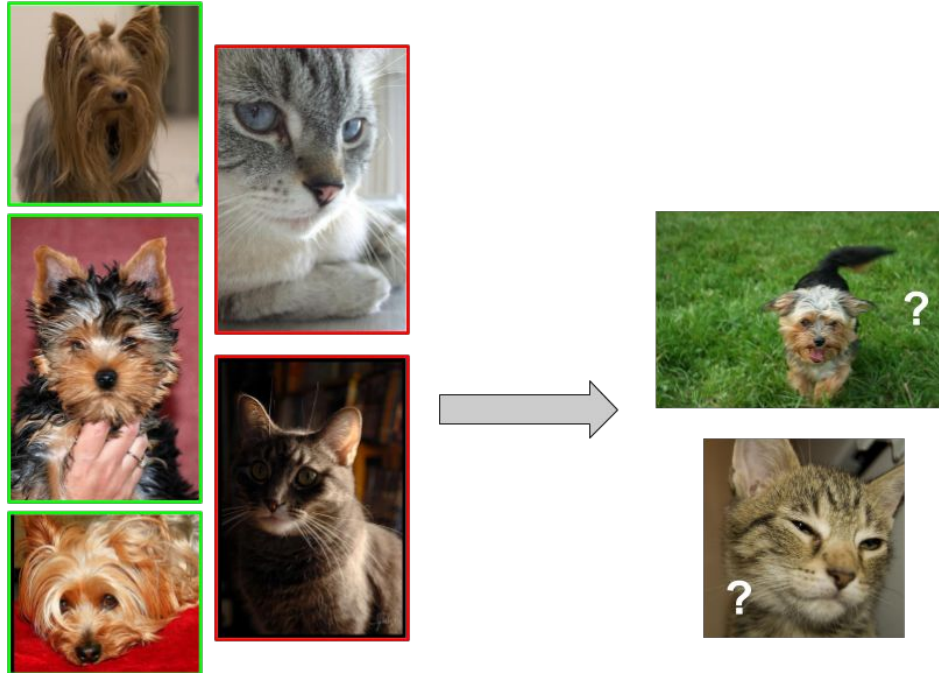
# Tips & Tricks



# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- **Tips & tricks**
- Exercises

# Generalization from small training set



# Overfitting curse

## Symptoms

- Very high training accuracy
- Very low testing accuracy

→ Model doesn't generalize to unseen data





# Regularization

## $L_2$ regularization

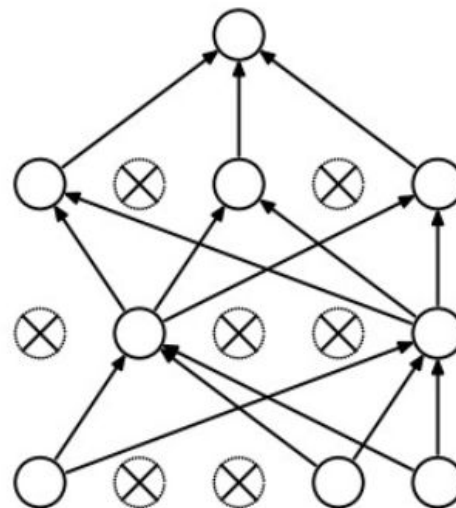
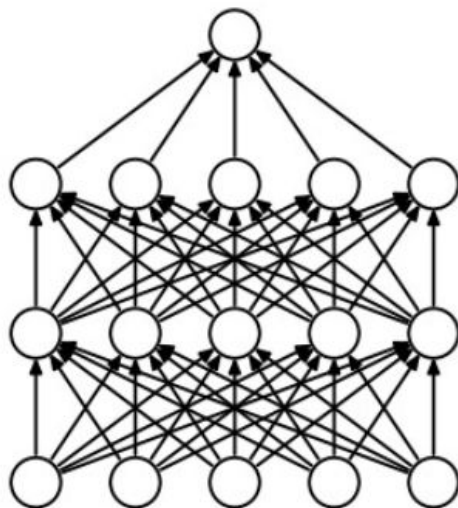
- Most common form of regularization
- Penalizing the squared magnitude of weights in the objective

## $L_1$ regularization

- Relatively common form of regularization
- Penalizing  $L_1$  of weights in the objective

# Dropout

Removing a neuron from a designated layer during training or by dropping certain connection







# Batch normalization

A common practice in NN, forces activations to have unit gaussian distribution

- Insert BN layer after FC and Conv, and before non-linearities

Robust networks to bad initialization

- interpreted as doing preprocessing at every layer of the network
- Normalization is differentiable



# Transfer learning

Problem Training a model from scratch only with small data is **challenging** and suffer from **overfitting**



# Transfer learning

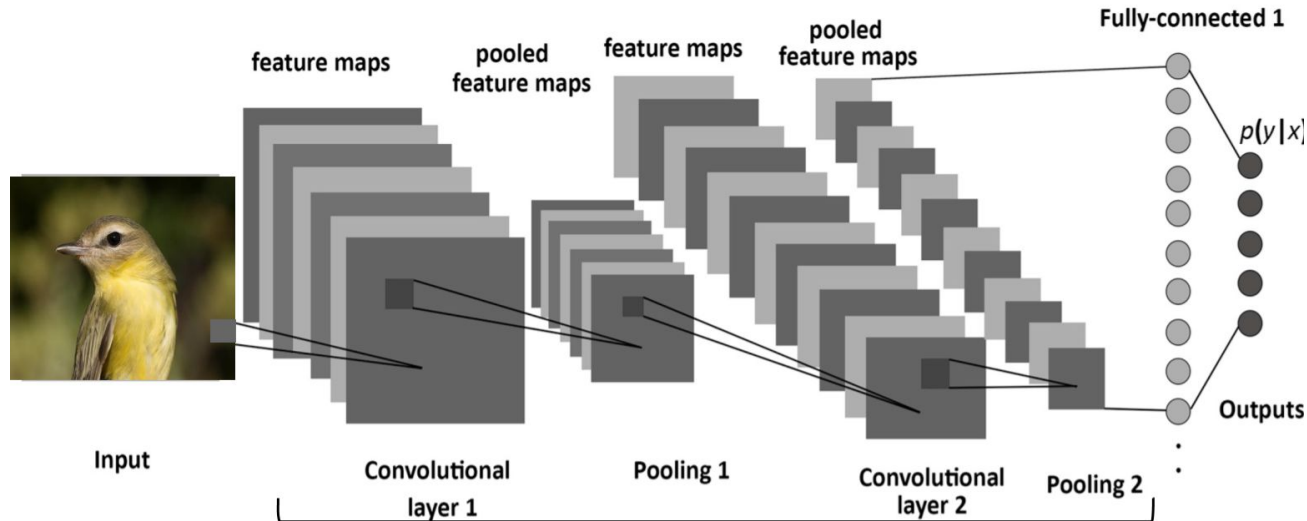
Problem Training a model from scratch only with small data is **challenging** and suffer from **overfitting**

Solution Use the knowledge of a pre-trained model on a border task to solve more specific one

- Bottle neck features
- Fine-tuning top layers

# Transfer learning

Extract bottleneck features from a pre-trained network





# Transfer learning

Finetune top layers of pre-trained network

- Remove top dense layers
- Add your own classification layers on top
- Freeze bottom layers
- Fine-tune top layers on small data

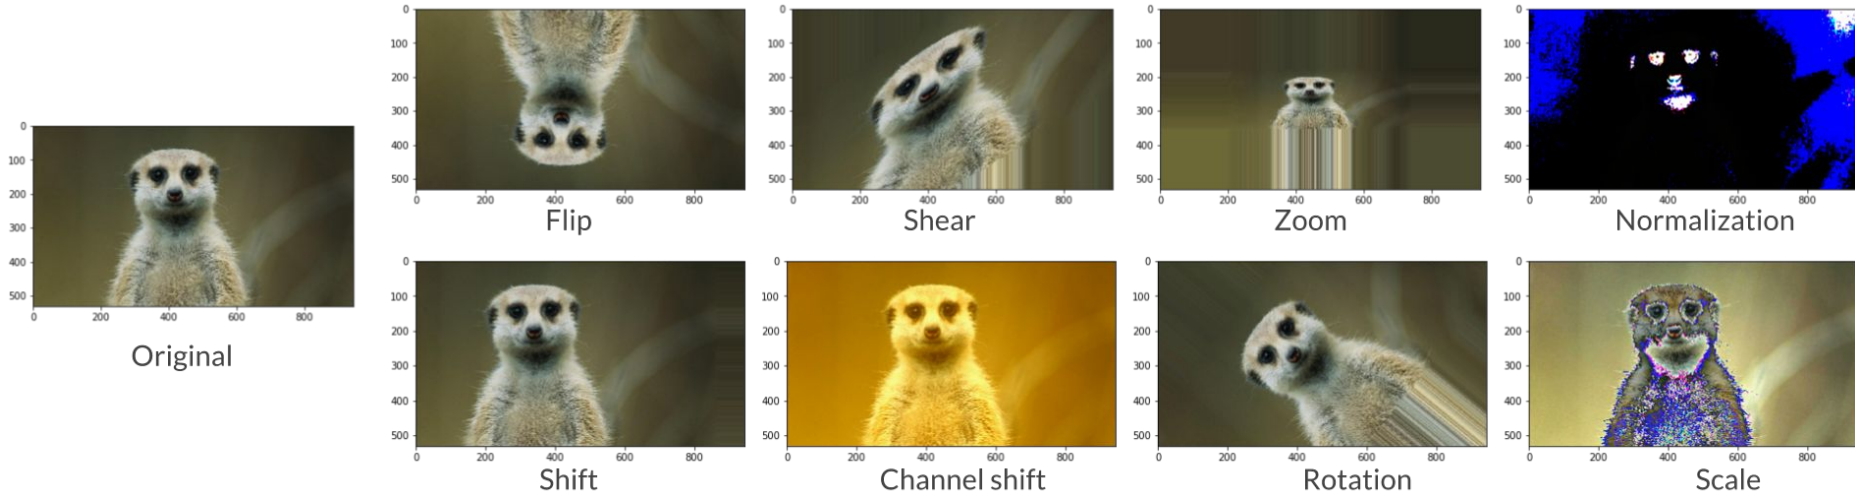


# Data augmentation

Increase the amount of training data using information only in our training data

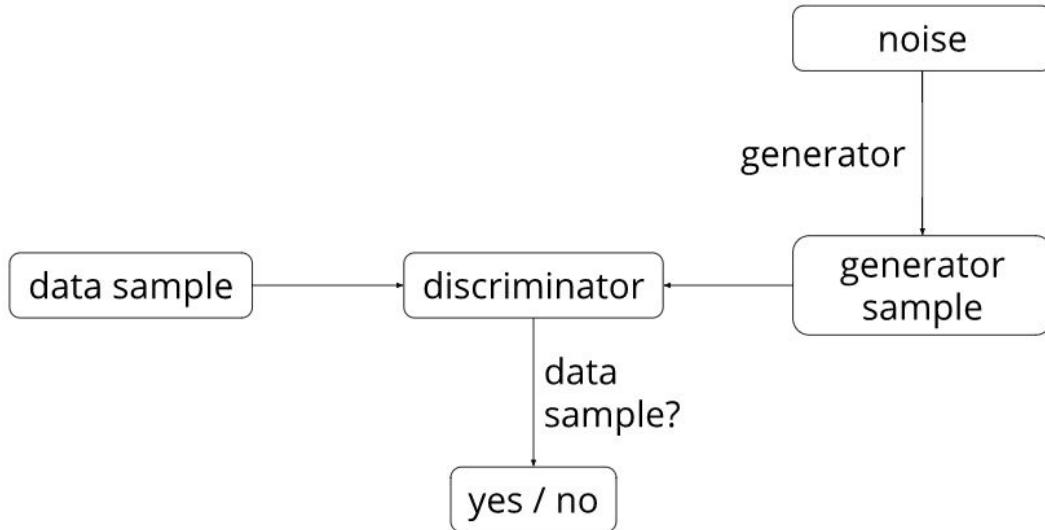
- Affine transformations
- Generative Adversarial Networks (**GANs**)

# Affine transformations



# Generative Adversarial Networks (GANs)

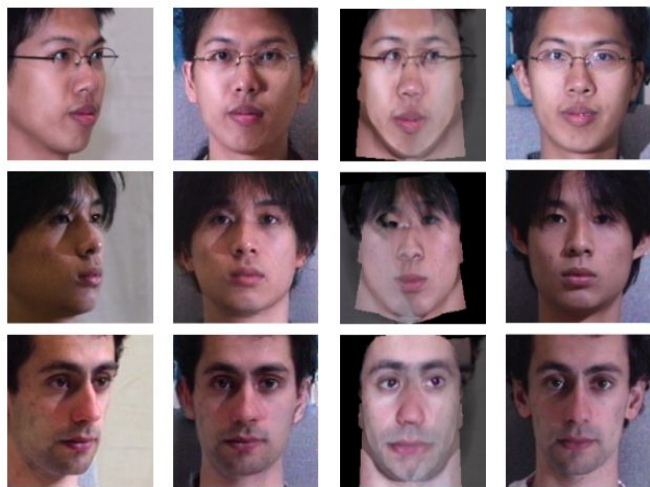
Learns the data distribution





# Generative Adversarial Networks (GANs)

Frontal face generator



GT profile

1

2

GT frontal

# Generative Adversarial Networks (GANs)

Image to image translation



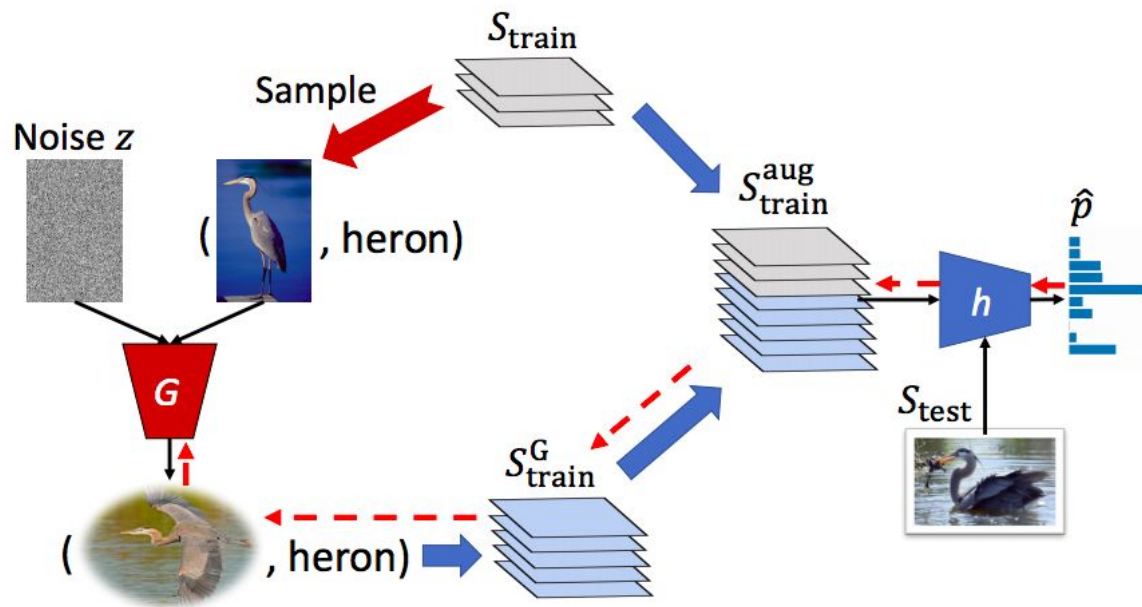
# Low-shot learning + GANs

Low-Shot Learning from Imaginary Data

Meta-learning + Hallucination



# Low-shot learning + GANs





# Q & A



# Structure

- Introduction & motivation
- Zero-shot learning
  - Definition
  - Side information
  - Zero-shot learning models
  - Exercise
- Low-shot learning
  - Definition
  - Low-shot learning models
- Tips & tricks
- Exercises



# **Image augmentation exercise**



# Image augmentation

- Clone git [repo](#)
- New notebook in notebooks/small\_classifier/
- Read and plot meerkat.jpg
- Use keras `keras.preprocessing.image.ImageDataGenerator` to generate different images
- Plot results





# **Image classifier with small set exercise**



# Image classifier with small set

Train small network from scratch

- Clone git [repo](#)
- New notebook in notebooks/small\_classifier/
- Define small conv net: 3 conv blocks (2Dconv,relu\_activation,max\_pooling) + 2 dense layers (don't forget flatten!)
- Compile network with binary\_crossentropy loss and rmsprop
- Define image generator
- Train using generator
- Bonus: plot loss & accuracy



# Image classifier with small set

Train small network from scratch

- Clone git [repo](#)
- Use image\_classifier\_with\_small\_data notebook in notebooks/small\_classifier/
- Define a model: 3 conv blocks (2Dconv,relu\_activation,max\_pooling) + 2 dense layers (don't forget flatten!)
- Compile network with binary\_crossentropy loss and rmsprop
- Define image generator
- Train using generator
- Bonus: plot loss & accuracy

## Cheatsheet

**Train image** 2000

**Validation images** 800

**Input size** 150x150x3 (w,h,RGB)

**Conv\_1:** filters 32, kernel size(3,3)

**Conv\_2:** filters 32, kernel size(3,3)

**Conv\_3:** filters 64, kernel size(3,3)

**Dense\_1:** 64, relu

**Dense\_2:** ? , sigmoid



# Image classifier with small set

Extract features from pre-trained model

- `from keras.applications import vgg16` (set `model_top` to `false`)
- Define image generator OR loop through images in data directory
- Use `model.predict_generator` to get features
- Save features in `.npy` file



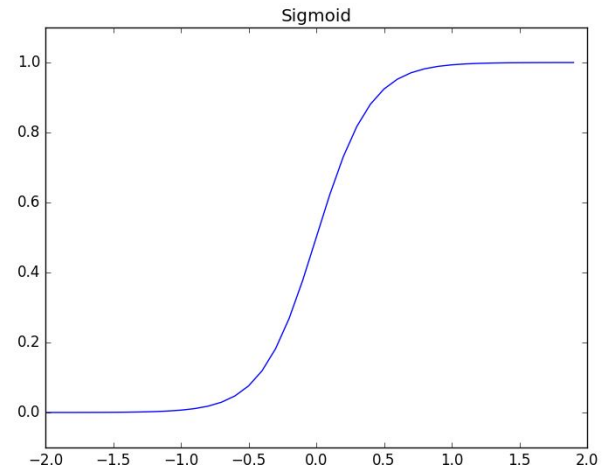
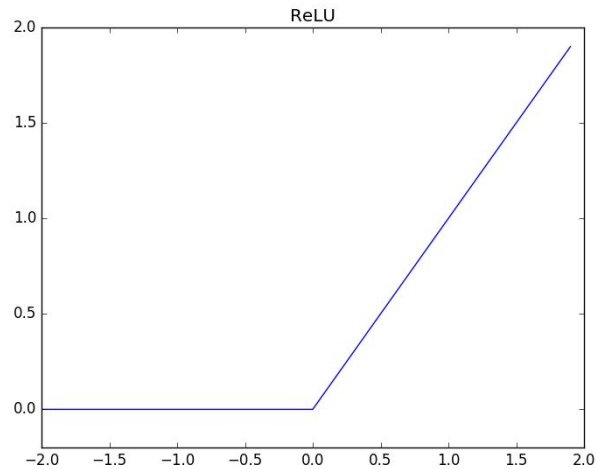
# Image classifier with small set

Train small MLP on bottleneck features

- Define network with two dense layers (don't forget activations and dropout)
- Compile with binary\_crossentropy loss and rmsprop
- Train with bottleneck features
- Bonus: plot loss & accuracy

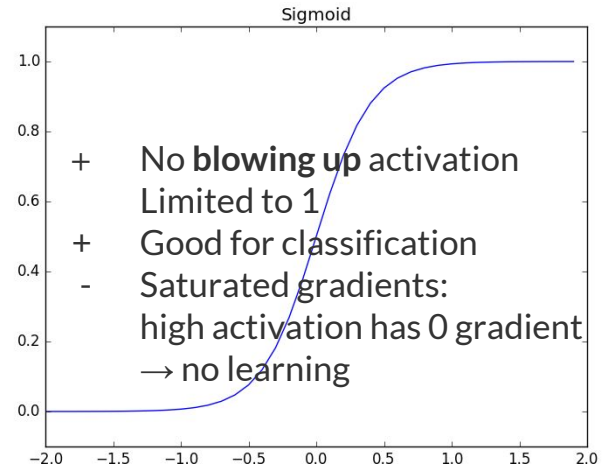
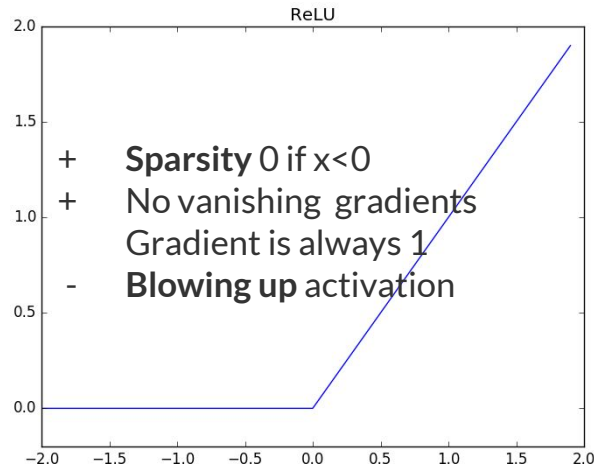
# Image classifier with small set

Relu vs. Sigmoid



# Image classifier with small set

Relu vs. Sigmoid





# Image classifier with small set

Train small MLP on bottleneck features

- Define network with two dense layers (don't forget activations and dropout)
- Compile with binary\_crossentropy loss and rmsprop
- Train with bottleneck features
- Bonus: plot loss & accuracy

## Cheatsheet

**Input size** image features size

**Flatten** 3D feature maps

**Dense\_1:** 256, relu

**Dropout:** rate 0.5

**Dense\_2:** 1, sigmoid





# Image classifier with small set

Fine-tune pre-trained network

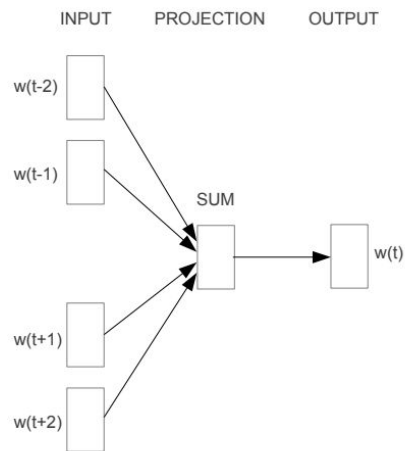
- Load pre-trained model vgg (weights same as before)  
Note: specify input size according to our images
- Create a new model (vgg + previous mlp)
- Freeze first 15 layers of the new model
- Compile new model with binary\_crossentropy loss and SGD with low learning rate (finetuning)
- Train with images



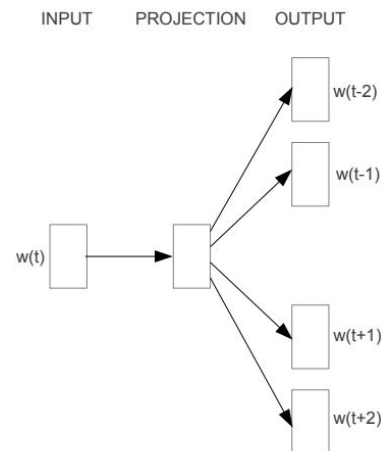
# Word Embedding

# Word embedding

Dense representations: Word2vec



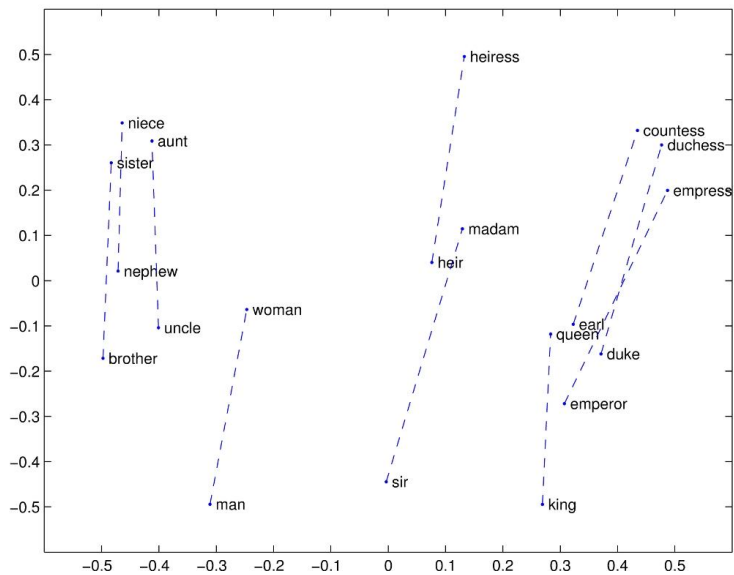
**CBOW**



**Skip-gram**

# Word embedding

Dense representations: Global vectors for word representation (Glove)





# Word embedding exercise



# Word embedding exercise

Train your own word2vec

- pip install gensim, tsne, bokeh
- Use sample corpus cloned from repo OR download sample text corpus eng\_news\_2005\_100K from <http://wortschatz.uni-leipzig.de/en/download/>
- Create new notebook train\_word2vec
- Train word2vec model using gensim
- Sanity check
- Tsne & plot with bokeh



# Word embedding exercise

Load pre-trained Glove

- Download pre-trained model from <https://nlp.stanford.edu/projects/glove/>
- Load the model using `gensim.models.KeyedVectors.load_word2vec_format`  
Note: fix first line format
- **Trick:** for quicker loading
- Sanity check
- Tsne & plot with bokeh



# Bonus





# Interviews QA

- Curse of Dimensionality



# Interviews QA

- **Curse of Dimensionality**
  - More features → harder to find a solution



# Interviews QA

- **Curse of Dimensionality**
  - More features → harder to find a solution
- **Bias-Variance Tradeoff**



# Interviews QA

- **Curse of Dimensionality**
  - More features → harder to find a solution
- **Bias-Variance Tradeoff**
  - Bias: error due to simplistic assumptions in the model, how well the model fits the data
  - Variance: error due to too much complexity in the model (sensitive for little changes), how much the model changes based on changes in the inputs



# Interviews QA

- Why Conv layer and not FC for images?



# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant



# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant
- **Max pooling?**



# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant
- **Max pooling?**
  - Reduce computation without losing too much information (max activation)





# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant
- **Max pooling?**
  - Reduce computation without losing too much information (max activation)
- **Normalization?**



# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant
- **Max pooling?**
  - Reduce computation without losing too much information (max activation)
- **Normalization?**
  - subtracting the mean of each data point and dividing by its standard deviation
  - makes all features weighted equally → stable convergence



# Interviews QA

- **Why Conv layer and not FC for images?**
  - Conv preserves spatial information in the image
  - Conv translation invariant
- **Max pooling?**
  - Reduce computation without losing too much information (max activation)
- **Normalization?**
  - subtracting the mean of each data point and dividing by its standard deviation
  - makes all features weighted equally



# Interviews QA

- Precision vs. Recall



# Interviews QA

- **Precision vs. Recall**
  - Recall: amount of positives your model claims compared to the actual number of positives
  - Precision: amount of correct positives your model claims compared to the number of positives it actually claims



# Interviews QA

- **Precision vs. Recall**
  - Recall: amount of positives your model claims compared to the actual number of positives
  - Precision: amount of correct positives your model claims compared to the number of positives it actually claims
- **F1 score**



# Interviews QA

- **Precision vs. Recall**
  - Recall: amount of positives your model claims compared to the actual number of positives
  - Precision: amount of correct positives your model claims compared to the number of positives it actually claims
- **F1 score**
  - weighted average of the precision and recall of a model
  - 1 the best, 0 the worst.
  - use it in classification where true negatives don't matter much.



## Extra resources

Some great resources if you want to dig deeper:

- [Blog post](#) about different optimization methods
- Matching networks for low-shot learning [code](#) implemented with tensorflow and [blog post](#)
- [Activation functions](#) comparison
- Stanford [course](#) on convolutional neural networks for visual recognition
- [Blog post](#) about using embedding layers in neural networks
- CVPR [tutorial](#) on zero-shot learning
- Interviews questions [springboard](#), [elitedatascience](#) and [towardsdatascience](#)





**Thanks!**